

Quantum-Inspired Cognition and Decision-Making for Autonomous Agents, from Robotics to AI/ML Systems: A Systematic Review

Yeshwanth Guru^{1*} and Dev Kunwar Singh Chauhan¹

^{1*}Department of Mechanical Engineering, Amrita Vishwa
Vidyapeetham, Chennai, India.

*Corresponding author(s). E-mail(s):

g_yeshwanth@ch.students.amrita.edu;

Contributing authors: c_devsingh@ch.amrita.edu;

Abstract

Autonomous robots and artificial intelligence (AI) systems increasingly decide under ambiguity and alongside people whose judgements violate classical probability. Quantum cognition models beliefs as vectors in a Hilbert space and judgements as measurements, so that question order, context and interference between unresolved possibilities change the resulting probabilities. These quantum-probability models run on ordinary computers and need no quantum hardware. This systematic review, reported following the PRISMA 2020 guideline, examines how such models have been transferred to autonomous agents. It identifies 67 application studies, grades each on a six-level evidence scale, appraises 61 against seven quality criteria and sets them against 71 theory and evidence studies of human judgement. The applications span quantum-like Bayesian networks, quantum-inspired machine learning, deep learning and audits of large language models, quantum-inspired reinforcement learning, robot perception and emotion models, multi-agent and human-machine teams, and human trust in AI. The psychological evidence is substantial but contested: order effects obeying the parameter-free quantum question equality are robust, whereas conjunction-fallacy and contextuality claims face credible classical rivals. The applications remain early: 24 of the 67 are proposals or formal models, 5 reach a physical system, only 24 of the 61 appraised studies compare fully against a classical model, and none shows in a controlled comparison that a quantum-probability module improves an autonomous system's task performance. The most defensible near-term use is human-facing: predicting human judgements,

trust and preferences. Case studies, an illustrative reinforcement-learning experiment, a validation roadmap and a benchmark suite with a replication checklist complete the review.

Keywords: Quantum cognition, Quantum probability, Decision-making, Human-robot interaction, Systematic review, Autonomous agents

Part A — Foundations (Sections 1–4)

1 Introduction — Why Quantum Cognition Matters for Robotics and AI

1.1 Motivation: Decision-Making Failures in Classical Robots and AI Systems

Classical autonomous systems rest on Kolmogorov probability. A robot’s belief is a probability distribution over states; new evidence is incorporated by Bayes’ rule; decisions maximise expected utility; and when the state is only partly observable the whole loop is formalised as a partially observable Markov decision process (Thrun et al. 2005; Lauri et al. 2023). Sensor fusion follows the same logic, with the Kalman filter as the optimal estimator for linear-Gaussian systems (Kalman 1960; Haykin 2001), and reinforcement learning (RL) assumes Markovian states and additive rewards (Sutton and Barto 2018). These assumptions have delivered reliable navigation, mapping and manipulation. They become fragile, however, in three situations that autonomous agents now meet routinely.

The first is interaction with people. Human judgement systematically violates classical probability: the probability of a conjunction is judged higher than that of one of its components (Tversky and Kahneman 1983), answers depend on the order in which questions are asked (Wang et al. 2014), and people who would choose the same option whatever the outcome of an uncertain event choose differently when the outcome is unknown, violating Savage’s sure-thing principle (Savage 1954; Shafir and Tversky 1992). A robot that predicts human intentions, calibrates trust or explains its choices with a model that cannot represent these patterns will be systematically wrong about its human partners (Song et al. 2022; Roeder et al. 2023). The second is ambiguity: when evidence is conflicting or the categories themselves are vague, a single joint probability distribution over all variables may not exist, a situation classical inference handles only with ad hoc devices (Ellsberg 1961; Blutner et al. 2013). The third is context and sequence: large language models (LLMs) and other learned systems also show order and framing effects (Binz and Schulz 2023; Tjautja et al. 2024), so AI components that query one another in sequence can inherit context-dependent behaviour that no fixed joint distribution describes (Lo et al. 2025).

1.2 The Quantum Cognition Hypothesis

Quantum cognition proposes that the mathematics of quantum theory, rather than its physics, provides a better description of these phenomena (Busemeyer and Bruza 2012; Pothos and Busemeyer 2013; Bruza et al. 2015). A belief state is a unit vector in a Hilbert space; a question is a subspace; the probability of an answer is the squared length of the state’s projection onto that subspace; and answering a question changes the state. Two principles do most of the work (Busemeyer and Wang 2015). Complementarity (incompatibility) means that some questions correspond to non-commuting projectors, so that answering one first changes the answer to the other and no joint distribution over both exists. Superposition means that, before a judgement is made, the state need not be in any definite answer; unresolved possibilities combine through amplitudes, and their interference can raise or lower probabilities relative to the classical law of total probability (Pothos and Busemeyer 2009). When all questions are compatible, quantum probability (QP) reduces exactly to classical probability, so the framework is a generalisation rather than a replacement (Bruza et al. 2015). The strongest single piece of support is a parameter-free prediction, the quantum question (QQ) equality, which held across 70 national surveys (Wang and Busemeyer 2013; Wang et al. 2014). The most recent overview of the research programme summarises two decades of such tests (Huang et al. 2025b).

1.3 Why Quantum Cognition is Viable Today (No Quantum Hardware Needed)

Quantum-cognition models are linear-algebra models: states are vectors or density matrices of modest dimension, measurements are projectors or instruments, and dynamics are unitary or Lindblad evolutions, all of which run on standard CPUs and GPUs. Quantum hardware is therefore optional. Circuits for cognitive decision models have been run on trapped-ion processors (Widdows et al. 2023) and compositional language models on superconducting devices (Lorenz et al. 2023), but these demonstrations show feasibility rather than advantage, and noisy intermediate-scale devices (Preskill 2018) add latency and error without adding modelling power at these sizes. A two- to eight-dimensional QP model can be evaluated in microseconds on an embedded processor (Section 20). The practical barrier to adoption is therefore not hardware but evidence: whether the effects are robust, whether QP models beat well-specified classical alternatives, and whether they improve what autonomous systems actually do.

1.4 Survey Scope: Robotics, ML, AI, DL, RL and Physical AI

We separate four kinds of model throughout: *quantum-like* models (QP mathematics on classical computers), *quantum-inspired* algorithms (classical heuristics borrowing quantum ideas, such as amplitude-based exploration), and quantum-circuit models that are either only designed or simulated, or executed on a quantum processing unit (QPU). The review covers (i) the psychological and mathematical foundations of quantum cognition, including open-system dynamics, quantum instruments and

contextuality; (ii) quantum-like graphical models for decision support; (iii) quantum-inspired machine learning (ML), deep learning (DL) and natural-language processing, including recent probes of LLMs; (iv) quantum-inspired RL and agent learning; (v) quantum-like models for robot perception, emotion, planning and embodied (physical) AI; (vi) multi-agent systems, swarms and human-machine teams; and (vii) quantum models of human trust in automation. Quantum computing used only to accelerate robot planning, learning, sensing or communication belongs to a companion survey (Guru and Chauhan 2026) and is not repeated here (Section 2.5.4). Claims that the brain is a quantum computer are treated only in Section 21.

The review is organised in five parts. Part A (Sections 1–4) sets out the motivation, the review method, the history and the mathematics. Part B (Sections 5–7) presents the theory and the deep and medium-depth treatment of perception, memory, decision-making, learning, planning and fusion. Part C (Sections 8–16) covers the applications to ML, DL, RL, embodied and multi-agent systems, positions them against prior work, compares them head to head with classical methods, and sets out a validation roadmap and case studies. Part D (Sections 17–20) discusses limitations, assumptions, ethics and implementation. Part E (Sections 21–25) connects the field to neuroscience, proposes benchmarks, compares related surveys, sets out open problems and concludes. Appendices A–E give derivations, code, raw results, benchmark protocols and extended case studies.

1.5 Survey Methodology

We report this review according to PRISMA 2020 (Page et al. 2021) and conducted it along the lines of the guidelines of Kitchenham and Charters (2007). The PRISMA 2020 checklist, with the location of each item, is in Online Resource 1 (Table S2). The protocol was not registered before the review began; a retrospective registration is deposited on the Open Science Framework [FILL: OSF registration DOI].

1.5.1 Research questions

The review addresses five questions. **RQ1**: which quantum-like, quantum-inspired and quantum-circuit models of cognition and decision have been applied to autonomous agents, and to which agent functions (perception and fusion, decision-making and planning, learning, embodied and social behaviour, multi-agent coordination, and human-AI trust)? **RQ2**: how strong is the psychological evidence for the QP phenomena on which these applications rely, and which classical explanations compete with it? **RQ3**: how far does the evidence for each application go, from conceptual proposal to field deployment, and how well do the studies meet basic quality criteria? **RQ4**: under which conditions do QP models outperform, match or underperform classical alternatives? **RQ5**: which validation programme, benchmarks and reporting standards are needed to move the field towards deployed agents?

1.5.2 Information sources and search

The corpus was assembled in four stages. *Stage 1*: a curated reading list of 102 items, compiled for this review and the companion survey, was checked item by item against

publisher, arXiv or repository records. About twenty entries needed corrections to titles, author lists, years or identifiers; three items that could not be located in any index were replaced by the two verified papers of the same research group; two items whose titles did not exist were replaced by the real works that matched their topic; and two further items were dropped, one because it could not be found and one because it did not concern quantum models. *Stage 2*: the PDFs held by the authors were identified by their first-page text rather than their file names, which added 31 relevant works. *Stage 3*: structured searches of Google Scholar, Semantic Scholar, arXiv, PubMed and publisher sites, combined with backward and forward citation chasing from the main reviews (Pothos and Busemeyer 2022; Huang et al. 2025b; Widdows et al. 2021; Meyer et al. 2022), ran in two streams: applications of quantum cognition to robots, agents, human-machine teams and reinforcement learning; and quantum-cognition theory, critiques, machine learning, LLMs and neuroscience published from 2016 to 2026. This added 74 verified works and gave a corpus of 205 works. *Stage 4*: a targeted search on 4 October 2026, with six topic queries and eleven verification queries, together with a cross-check against the companion survey’s corpus, added four application studies. The exact strings of stage 3 and the number of records each returned were not logged, so they cannot be reported; Online Resource 1 (Table S6) gives the concept blocks that reproduce its scope and the exact queries, dates and decisions of stage 4. The corpus was closed on 4 October 2026; its most recent study appeared in July 2026.

1.5.3 Eligibility criteria

A work entered the corpus if it (I1) proposed, evaluated or reviewed a quantum-like, quantum-inspired or quantum-circuit model of cognition, decision or learning applied to an artificial agent, a machine-learning or AI system, or human-AI or human-robot interaction; (I2) reported theory, empirical tests or critiques of QP models of human judgement and decision on which such applications rely; or (I3) supplied reviews, classical baselines, hardware, software or background needed for comparison. A work was excluded if it (E1) used the word “quantum” only as a metaphor; (E2) was written in a language other than English; (E3) repeated an included work (we kept the most complete version); (E4) concerned physical quantum processes in the brain without a cognitive model (these are discussed, ungraded, in Section 21); (E5) used quantum computation only to accelerate a robot or learning task, such as optimisation, search, variational-circuit reinforcement learning, sensing or communication, without a cognitive or decision model; or (E6) was not published in a journal, conference, book, thesis or preprint server (for example, project web pages). Works meeting I1 are the *application studies* (the primary studies of this review); works meeting only I2 are *theory and evidence studies*; and works meeting only I3 are reviews and background works, which were not graded. Eleven studies of the earlier corpus fell under E5 and are reviewed in the companion survey (Guru and Chauhan 2026): variational-circuit reinforcement learning (Chen et al. 2020; Jerbi et al. 2021; Skolik et al. 2022; Acuto et al. 2022; Sinha et al. 2025; Hohenfeld et al. 2024), quantum multi-agent reinforcement learning (Yun et al. 2022), quantum speed-ups of agent deliberation (Paparo et al. 2014), Grover-based motion planning (Chella et al. 2022), quantum-circuit swarm rules (Mannone et al. 2022) and quantum communication between cooperating robots (Khoshnoud

et al. 2020). They are listed with the reason for exclusion in Online Resource 1 (Table S7).

1.5.4 Selection process

Records were screened on title and abstract and then, where a full text could be obtained, in full text. The first author screened the records and extracted the data, and the second author reviewed the inclusion decisions, evidence levels and appraisal judgements; disagreements were resolved by discussion **[FILL: authors to confirm this sentence]**. Because the review was not independent double screening, inter-rater agreement could not be measured. Figure S1 in Online Resource 1 shows the flow of records with the counts that were documented. The final corpus holds 212 works: 67 application studies, 71 theory and evidence studies, 17 reviews, 43 background works, the 11 studies excluded under E5, which are cited only to mark the boundary with the companion survey, and three methodological references.

1.5.5 Data extraction and quality appraisal

For each application study we extracted the domain, the task or phenomenon, the model type, the method, the setup and data, the key result, the main limitation and the publication status (Online Resource 1, Table S1, and Online Resource 3), and wrote a structured summary (Online Resource 1, Section S9). Full texts could be read for 16 of the 67 application studies and 21 of the 71 theory and evidence studies; the other studies were extracted from their abstracts and publisher pages, and the basis of each extraction is recorded in the dataset. Where only an abstract was available and the level was unclear, the study received the lower level.

Application studies above level E1 were appraised against seven criteria (Table S4): Q1 model specification (state space, operators and free parameters stated); Q2 classical comparison (a classical Bayesian, Markov or neural model evaluated on the same task); Q3 empirical grounding (human data or a realistic agent or AI task rather than a toy example); Q4 parameter discipline (parameters predicted or constrained a priori, or the model validated on held-out data); Q5 quantitative evaluation (metrics with some measure of variability); Q6 reproducibility (code or data released); and Q7 critical discussion (limitations and alternative explanations). Each criterion received one of three judgements: met, partly met or not met. We report the judgements criterion by criterion rather than as a total, because a total would conceal where a study is weak; Table S5 and Online Resource 4 give them for every study. Reproducibility (Q6) can only be judged from the full text, so it was left unassessed for the 47 appraised studies read from their abstracts.

1.5.6 Classification scheme and threats to validity

Evidence levels

Each application and theory study was graded on the six levels of Table 1. The levels grade how far the evidence goes towards an autonomous system, not the quality of the underlying science: a rigorous psychology experiment is E3 because it involves

human data but no agent. The companion survey grades robotics studies on a different six-level scale, which separates simulation, QPU execution and robot-in-the-loop operation (Guru and Chauhan 2026). That scale has no rung for human data, which is the first test of a cognitive model, so it is not used here.

Model type and domain

Each application study was tagged as quantum-like, quantum-inspired, quantum circuit (design or simulation) or quantum circuit (QPU), and assigned to one of seven domains: decision models and networks; machine learning, natural-language processing and information retrieval; deep learning and LLMs; reinforcement learning and agent learning; robots and embodied agents; multi-agent systems and teams; and human-AI trust.

Table 1 Evidence levels used to grade application and theory studies, with the number of application studies (n = 67) and of all graded studies (n = 138) at each level

Level	Label	Definition	Typical example	n (appl.)	n (all)
E1	Conceptual	Position or proposal without a formal model or data	Architectural proposals (Aerts et al. 2011b; Lawless 2020)	6	8
E2	Formal model	Mathematical model or algorithm; analytic or toy illustration only	Quantum emotion spaces (Yan et al. 2019)	18	49
E3	Human data	Model fitted to or tested on human behavioural or neural data	QQ equality in surveys (Wang et al. 2014); trust in AI (Roeder et al. 2023)	9	46
E4	Simulation or benchmark	Simulated agents, ML benchmark datasets or LLM probes	Quantum-like LLM contextuality (Lo et al. 2025)	29	30
E5	Physical system	Physical robot or real hardware system in the laboratory	Quantum-inspired RL on a mobile robot (Dong et al. 2012)	5	5
E6	Field deployment	Real-world deployment or large-scale study with real users or operators	None found	0	0

Threats to validity

Coverage: the terminology is inconsistent (“quantum-like”, “quantum-inspired” and “quantum” are used interchangeably), several recent works are preprints that have not been peer-reviewed, and non-English work is missing; the four studies found by the October 2026 targeted search show that the earlier searches were not exhaustive.

Reporting: the strings and record counts of the main searches were not logged, so search completeness cannot be measured and the flow diagram reports only documented counts.

Extraction: 51 of the 67 application studies were extracted from their abstracts rather than their full texts, which may have led to errors in the assigned level and leaves reproducibility unassessed for most studies; ambiguous cases received the lower level.

Selection and grading: a single author screened the records and extracted the

data, and the second author reviewed rather than independently repeated the grading, so grading reliability cannot be measured; every judgement is published so that readers can check it. *Synthesis*: the heterogeneity of tasks and metrics prevented a quantitative meta-analysis, so the synthesis is narrative. *Publication bias* probably favours positive results.

2 Evolution of Quantum Cognition: From Psychology to AI and Robotics

Figure 1 summarises five eras in the development of quantum cognition and its transfer to autonomous agents.




Fig. 1 Five eras of quantum cognition and its transfer to autonomous agents, with representative works

2.1 Era 1 (1925–2000): Classical Cognitive Foundations and Early Quantum Ideas

2.1.1 Von Neumann and Dirac: Quantum Formalism Emerges (1925–1932)

The formalism that quantum cognition borrows was consolidated by von Neumann, who represented states as vectors in a Hilbert space, observables as self-adjoint operators and measurement as projection ([von Neumann 1932](#)). Two of his results matter here. First, quantum probability is defined on a lattice of subspaces that is not distributive, so the classical identity $P(A) = P(A \wedge B) + P(A \wedge \neg B)$ can fail. Second, measurement changes the state, the property later exploited to model the constructive nature of human judgement.

2.1.2 Classical Cognitive Revolution (1950s–1980s): Bayesian Networks and Symbols

Savage’s axiomatisation of subjective expected utility made the sure-thing principle a pillar of rational choice ([Savage 1954](#)). Cognitive science then developed along two lines that still shape autonomous systems: symbolic architectures, from which today’s cognitive architectures descend ([Kotseruba and Tsotsos 2020](#); [Laird et al. 2017](#)), and probabilistic models, later formalised as Bayesian networks and probabilistic robotics ([Thrun et al. 2005](#)). Both assume that all variables can be placed in one sample space.

2.1.3 Tversky and Kahneman: Decision Anomalies Surface (1974–2000)

The heuristics-and-biases programme documented systematic departures from probability theory ([Tversky and Kahneman 1974](#)). The conjunction fallacy (judging “bank teller and feminist” more probable than “bank teller”) ([Tversky and Kahneman 1983](#)), ambiguity aversion in the Ellsberg paradox ([Ellsberg 1961](#)) and the disjunction effect in the Prisoner’s Dilemma ([Shafir and Tversky 1992](#)) became the benchmark phenomena that quantum models would later address.

2.2 Era 2 (2000–2010): Quantum Probability Enters Cognitive Science

2.2.1 Aerts and Gabora: First Quantum Cognition Papers

The Brussels group argued that concepts behave like quantum entities whose meaning is actualised by context. Aerts, D’Hondt and Gabora showed that the quantum-logical disjunction is not classical ([Aerts et al. 2000](#)), and Aerts later showed that membership judgements for combined concepts (the “guppy effect” and over- and underextension) are incompatible with classical models but reproduced in Fock space ([Aerts 2009](#)).

2.2.2 Khrennikov: Quantum-Like Formalism in Cognition

Khrennikov developed quantum-like models of information spaces ([Khrennikov 1999](#)) and a contextual probability framework in which interference arises from incompatible contexts rather than physical quantum effects ([Khrennikov 2010](#)). The “quantum-like” label used throughout this review originates in this work.

2.2.3 Early Applications: Prisoner’s Dilemma, Concept Combinations

Busemeyer, Wang and Townsend introduced quantum dynamics of decision-making, contrasting quantum and Markov random walks ([Busemeyer et al. 2006](#)). Pothos and Busemeyer then showed that interference between the “opponent defects” and “opponent cooperates” branches reproduces the disjunction effect that classical models cannot ([Pothos and Busemeyer 2009](#)), and Busemeyer, Wang and Lambert-Mogiliansky derived parameter-free empirical tests separating Markov and quantum models ([Busemeyer et al. 2009](#)). Bruza and colleagues asked whether the human mental lexicon is quantum-like ([Bruza et al. 2009](#)).

2.3 Era 3 (2011–2015): The Busemeyer and Bruza Framework Emerges

2.3.1 Busemeyer and Bruza, “Quantum Models of Cognition and Decision” (2012)

The monograph by Busemeyer and Bruza unified the field’s principles, models and data ([Busemeyer and Bruza 2012](#)); it remains the standard reference (a second edition has since appeared). Its companion target article in *Behavioral and Brain Sciences*, with open peer commentary, made the case to mainstream cognitive science ([Pothos and Busemeyer 2013](#)).

2.3.2 Order Effects and Conjunction Fallacy Formalised

Busemeyer, Pothos, Franco and Trueblood explained conjunction and disjunction errors as sequential projections ([Busemeyer et al. 2011](#)); Trueblood and Busemeyer modelled order effects in inference ([Trueblood and Busemeyer 2011](#)); and Wang and Busemeyer derived the QQ equality ([Wang and Busemeyer 2013](#)), confirmed in 70 national surveys ([Wang et al. 2014](#)).

2.3.3 Interference Effects: Two-Stage Measurements in Decisions

Two-stage paradigms, in which an intermediate judgement is either elicited or not, became the main test of interference. Making a categorisation before an action ([Wang and Busemeyer 2016](#)), or a choice before a confidence rating ([Kvam et al. 2015](#)), changed later judgements in ways that violate the law of total probability.

2.3.4 Widespread Adoption in Psychology (Special Issues, Conferences)

A special issue of *Topics in Cognitive Science* set out the case for quantum models of cognition (Wang et al. 2013) and extended them to similarity (Pothos et al. 2013), vagueness (Blutner et al. 2013) and concept dynamics (Aerts et al. 2013). The Quantum Interaction conference series provided a venue shared with information retrieval and AI, where the first explicit bridge to AI and robotics was proposed (Aerts et al. 2011b,c). Survey articles consolidated the field (Ashtiani and Azgomi 2015; Bruza et al. 2015).

2.4 Era 4 (2016–2020): Beyond Psychology — First Robotics and AI Connections

2.4.1 Quantum Cognition Applied to Autonomous Systems and ML

Quantum-like Bayesian networks replaced classical probabilities with complex amplitudes and gave decision-support systems a way to represent sure-thing violations (Moreira and Wichert 2016; Moreira 2017). In information retrieval and NLP, density-matrix language models (Sordoni et al. 2013) led to quantum-inspired neural networks for matching and sentiment (Zhang et al. 2018; Li et al. 2019). In RL, a separate line of quantum-inspired algorithms had already been validated on a mobile robot (Dong et al. 2008, 2012), and projective simulation offered a physics-motivated learning agent (Briegel and De las Cuevas 2012) tested on real robotic manipulation (Hangl et al. 2016). Order effects were proposed as a design concern for the queries of intelligent agents (Kak 2018).

2.4.2 Exploration of Sensor Fusion and Perception Problems

Quantum-inspired multimodal fusion treated modalities as incompatible measurement contexts (Li et al. 2021a; Gkoumas et al. 2021). For robots, Lanza, Solinas and Mastrogiovanni proposed a quantum-like perception model in which percepts are qubit states and queries are measurements (Lanza et al. 2020), extended to multi-sensory integration (Lanza et al. 2021).

2.4.3 Hybrid Quantum-Classical Architectures for Real-Time Control

Early “quantum robot” architectures combined classical controllers with (then hypothetical) quantum processors (Benioff 1998; Dong et al. 2006); robot emotion was modelled as a qubit-encoded emotion space (Yan et al. 2019) and as mixed states of reflection and affect (Hoorn and Ho 2019); and a quantum-like cognitive architecture was proposed as an alternative to resource-rational analysis (Moreira et al. 2020a). None of these architectures closed a real-time control loop with a quantum processor.

2.5 Era 5 (2021–Present): Quantum Cognition for Autonomous Systems

2.5.1 Fault Diagnosis and Multi-Goal Trade-offs in Autonomous Robots

Recent robot-facing work includes quantum-like pedestrian-intention models for automated driving (Song et al. 2022), quantum affective decision models for social robots (Ho and Hoorn 2022; Yan et al. 2021c), quantum-circuit teleo-reactive controllers (Chella et al. 2026), quantum game models of driving interactions (Essalmi et al. 2025, 2026) and a social-robotics research programme (De Carolis et al. 2025). Multi-goal trade-offs appear in quantum-like evidence networks (Gao et al. 2025) and multi-attribute decision models (She et al. 2021). We found no study that applies quantum cognition to fault diagnosis in robots; the closest work is a fault-tolerant quantum filter for quantum systems (Gao et al. 2021), which addresses a different problem.

2.5.2 ML and AI Integration: Feature Learning and Decision-Making

Quantum-cognitive fusion (Gkoumas et al. 2021), density-matrix matching (Zhang et al. 2025) and quantum-cognitive neural networks that reproduce human-like confidence (Maksimovic and Maksymov 2025) extend quantum-inspired ML, and complex-valued sequence models with Born-rule decoding carry the same ideas into language modelling (Nebli et al. 2026). Since 2024, a new line probes LLMs for quantum-like structure: contextuality in LLM-derived probabilities (Lo et al. 2025), context effects in similarity judgements (Upreti et al. 2024), Bell-type correlations in concept combinations (Aerts et al. 2026), and probability-judgement fallacies (Imannezhad et al. 2026; Kang 2026).

2.5.3 RL Applications: Policy Optimisation Under Uncertainty

Quantum-inspired experience replay continued the classical quantum-inspired RL line (Wei et al. 2022), while variational-circuit RL grew into a separate field that the companion survey reviews (Guru and Chauhan 2026; Meyer et al. 2022). On the cognitive side, open-system quantum–Markov models describe how preferences and beliefs evolve (Busemeyer et al. 2020; Kvam et al. 2021; Epping et al. 2023), and quantum models of trust in AI track how human reliance changes during interaction (Roeder et al. 2023; van der Meer et al. 2025; Humr et al. 2025). Both are candidate models of the human reward-giver in RL from human feedback (Section 10.1).

2.5.4 Integration with Survey 1 (Quantum Robotics Comprehensive Overview)

This review is the second of a pair. The companion systematic survey (Guru and Chauhan 2026) covers quantum computing, sensing and communication for robots: quantum optimisation and search for planning and scheduling, variational-circuit reinforcement learning, quantum navigation sensors and quantum communication between robots, graded by how close each result comes to a working robot. It found that no

quantum-computing method has yet reached deployment on a robot, whereas quantum sensing and communication have, and it treats quantum cognition in a single short section. The present review is the in-depth treatment of the cognitive side. Quantum cognition needs no quantum processor, so its bottleneck is empirical validation rather than hardware. Where the two meet, for example quantum RL and circuit implementations of cognitive models, we refer to the companion survey for hardware constraints. Studies that use quantum computation only as an accelerator are excluded here (criterion E5) and cited only to mark the boundary; eight application studies are discussed in both reviews, from different angles, and are flagged in the dataset (Online Resource 3).

3 Mathematical Preliminaries — Linear Algebra and Hilbert Spaces (for AI/ML/Roboticians)

This section gives the minimum machinery needed to read the rest of the review; derivations are collected in Appendix A, and tutorials are available elsewhere (Yearsley and Busemeyer 2016; Yearsley 2017).

3.1 Complex Vector Spaces and Inner Products

A state is a vector $|\psi\rangle$ in $\mathbb{C}^n$. The inner product $\langle\phi|\psi\rangle = \sum_i \phi_i^* \psi_i$ is linear in its second argument and conjugate-linear in its first; the norm is $\|\psi\| = \sqrt{\langle\psi|\psi\rangle}$. Complex amplitudes matter because probabilities are squared magnitudes: two paths to the same outcome add as amplitudes, and their relative phase determines whether they reinforce or cancel.

3.2 Hilbert Spaces and Orthonormality

A Hilbert space is a complete inner-product space; all models in this review are finite-dimensional, so it is simply $\mathbb{C}^n$ with the standard inner product. An orthonormal basis $\{|e_i\rangle\}$ satisfies $\langle e_i|e_j\rangle = \delta_{ij}$. In cognitive models each basis vector represents a fully resolved answer pattern, and different questions correspond to different bases of the same space.

3.3 Linear Operators and Matrices

Observables are Hermitian operators ($A = A^\dagger$); dynamics are unitary operators ($U^\dagger U = I$) that preserve norms, generated by a Hermitian Hamiltonian H through $U(t) = \exp(-iHt)$. Two operators are compatible when they commute ($AB = BA$). The commutator $[A, B] = AB - BA$ measures incompatibility and is the formal source of order effects.

3.4 Eigenvalues, Eigenvectors and Projectors

A Hermitian operator has real eigenvalues and an orthonormal basis of eigenvectors; each distinct answer to a question corresponds to an eigenvalue, and the eigenvectors with that eigenvalue span the answer’s subspace. A projector P satisfies $P = P^\dagger =$

P^2 ; it projects onto that subspace. The Born rule gives $p(\text{answer}) = \|P|\psi\rangle\|^2 = \langle\psi|P|\psi\rangle$, and the Lüders rule gives the post-answer state $P|\psi\rangle/\|P|\psi\rangle\|$. Sequential judgements therefore have probability $p(A \text{ then } B) = \|P_B P_A |\psi\rangle\|^2$, which differs from $p(B \text{ then } A)$ when P_A and P_B do not commute. Mixed states are density matrices ρ (positive, unit trace), with $p = \text{tr}(P\rho)$; general measurements are positive-operator-valued measures (POVMs) or quantum instruments, which are needed to model some combinations of effects (Section 5.2) (Ozawa and Khrennikov 2020; Fuyama et al. 2025).

3.5 Tensor Products and Multi-Dimensional Spaces

Composite systems live in tensor products: two binary questions that are compatible can be represented in $\mathbb{C}^2 \otimes \mathbb{C}^2 = \mathbb{C}^4$. States that cannot be written as products (entangled states) represent dependencies that are not reducible to separate beliefs, and have been used to model concept combination (Aerts and Sozzo 2011). For ML practitioners, tensor products are also the basis of tensor-network models (Stoudenmire and Schwab 2016). The dimension grows as 2^n for n binary variables, which is the central computational limitation of naive QP models (Section 20). Hilbert-space multidimensional theory addresses this by choosing low-dimensional representations tailored to the questions asked (Busemeyer and Wang 2018).

4 Standardised Terminology — Glossary and Visual Diagrams

4.1 Key Terms: Superposition, Entanglement, Interference, Measurement

In quantum cognition these terms are used in a precise mathematical sense that should not be confused with claims about physical quantum states in the brain. Superposition denotes a state that is a linear combination of answer states, representing unresolved uncertainty that is not mere ignorance of a fixed answer. Measurement is the act of eliciting a judgement or decision, which projects the state and so changes subsequent judgements. Interference is the cross term that appears when the probability of an outcome is computed from two unresolved paths; it is responsible for violations of the law of total probability. Entanglement is non-separability of a composite state; in cognition it is usually diagnosed by Bell-type or CHSH inequalities, and such diagnoses require care because signalling between measurements (direct influence of context on marginal probabilities) can mimic it (Dzhafarov and Kujala 2016).

4.2 Quantum vs. Classical Probability

Table 2 compares the two frameworks. Classical probability assigns probabilities to subsets of one sample space and is commutative and distributive; QP assigns them to subspaces, is non-commutative for incompatible questions and violates the distributive law. Both satisfy the basic axioms for a single set of compatible events, and QP reduces to classical probability when all projectors commute (Bruza et al. 2015). The difference

matters only for incompatible questions, which is why careful experimental design is needed to discriminate the two.

Table 2 Classical versus quantum probability

Aspect	Classical (Kolmogorov) probability	Quantum probability
Events	Subsets of one sample space	Subspaces of a Hilbert space
State	Probability distribution	Unit vector $ \psi\rangle$ or density matrix ρ
Probability rule	Measure of the subset	Born rule: $\ P \psi\rangle\ ^2$ or $\text{tr}(P\rho)$
Updating	Bayes' rule (conditioning)	Lüders rule (projection) or quantum instrument
Conjunction	Commutative	Sequential and order-dependent for incompatible questions
Law of total probability	Always holds	Can be violated by interference terms
Distributive law	Holds	Fails for incompatible events
Joint distribution	Always exists	Exists only for compatible questions
Dynamics	Markov (Kolmogorov) equations	Unitary (Schrödinger) or open-system (Lindblad) equations
Relation	–	Reduces to classical probability when all projectors commute

4.3 Visual Diagrams: Bloch Sphere, State Vectors, Measurement Bases

Figure 2 illustrates the geometry of an order effect and Figure 3 the standard geometric picture for a single binary question; interference is illustrated in Figure 4 (Section 5.5). On the Bloch sphere, a pure state of one qubit is a point on the surface; a binary question is a diameter; the probability of “yes” is $\cos^2(\theta/2)$, where θ is the angle between the state and the “yes” pole. Two questions whose diameters are not aligned are incompatible.

4.4 Glossary of Symbols and Notation

Table 3 lists the symbols and abbreviations used in the review.

Part B — Theory (Sections 5–7)

5 Foundational Theory — Busemeyer, Bruza, Order Effects, Conjunction Fallacy

5.1 Overview of Quantum Cognition as a Psychological Theory

Quantum cognition is best understood as a family of formal models, not a single theory of mind. Its common core is the use of QP to represent judgement and decision processes in which the act of judging is constructive: a person does not read out a pre-existing answer but creates one in the context of the question, and that





Fig. 2 Geometry of a question-order effect in a real two-dimensional belief space. The state $|\psi\rangle$ is projected onto “A = yes” then “B = yes” (blue path) or in the reverse order (orange path). With the angles used in Example 1 (Section 5.3), $p(A \text{ then } B) = 0.29$ and $p(B \text{ then } A) = 0.22$, while the QQ equality is satisfied exactly

answer becomes part of the context for the next question (Busemeyer and Bruza 2012; Pothos and Busemeyer 2013). On this view the probabilistic “errors” catalogued by the heuristics-and-biases programme are not failures of a classical reasoner but signatures of a different, and arguably principled, probability calculus (Pothos et al. 2017). The Annual Review by Pothos and Busemeyer (Pothos and Busemeyer 2022) and the 2025 overview of the research programme (Huang et al. 2025b) are the authoritative summaries; Busemeyer and Wang give an accessible introduction (Busemeyer and Wang 2015).

5.2 The Quantum Cognition Framework: Core Postulates

Following Busemeyer and Bruza (2012) and Yearsley and Busemeyer (2016), a quantum cognitive model specifies five elements:

- A Hilbert space whose dimension is set by the number of distinguishable answer patterns (often 2–4 dimensions per question set).
- An initial belief state $|\psi\rangle$ (or density matrix ρ) before any question is asked.
- A projector (or, more generally, an instrument) for each answer to each question, with incompatible questions represented in different bases.
- Born and Lüders rules for the probability of an answer and the updated state.
- Optionally, a Hamiltonian or Lindblad generator for how the state evolves between measurements, for example as evidence accumulates (Busemeyer et al. 2006, 2020).



$$P(\text{yes to A}) = \cos^2(\theta/2); \text{ A and B are complementary bases (non-commuting)}$$

Fig. 3 Bloch-sphere representation of a single binary judgement. Complementary questions A and B correspond to orthogonal axes; the state $|\psi\rangle$ has definite probabilities for both, but no joint distribution over the two answers

Table 3 Glossary of symbols and notation

Symbol	Meaning
$ \psi\rangle, \langle\psi $	State vector (ket) and its conjugate transpose (bra)
$\langle\phi \psi\rangle$	Inner product
ρ	Density matrix (mixed state); $\text{tr } \rho = 1, \rho \succeq 0$
$P_A, I - P_A$	Projectors onto the “yes” and “no” subspaces of question A
$[A, B] = AB - BA$	Commutator; zero for compatible questions
$\ P \psi\rangle\ ^2$	Born-rule probability
$U(t) = \exp(-iHt)$	Unitary evolution generated by Hamiltonian H
L_k, γ_k	Lindblad (jump) operators and rates in open-system models
$\otimes$	Tensor product of spaces or states
θ	Interference phase between two branches
QQ	Quantum question equality
QLBN	Quantum-like Bayesian network
QRL, QiRL	Quantum and quantum-inspired reinforcement learning
PQC, VQC	Parametrised and variational quantum circuit
PS	Projective simulation
CbD	Contextuality-by-default
POVM	Positive-operator-valued measure (generalised measurement)
QPU	Quantum processing unit

Two refinements are now standard. First, projective measurement alone cannot reproduce both question-order effects and response replicability (repeating a question returns the same answer), a result known as the response-replicability problem (Khrennikov et al. 2014). Quantum instruments, more general state-update maps, resolve this (Ozawa and Khrennikov 2020, 2021), and recent work argues that the non-classicality of cognition lies in non-commuting state updates rather than non-commuting observables (Fuyama et al. 2025; Busemeyer et al. 2025). Second, pure unitary dynamics predicts perpetual oscillation, so open-system (quantum–Markov, Gorini–Kossakowski–Sudarshan–Lindblad) models, in which the belief state couples

to an environment and settles into a decision, have become the preferred dynamical framework (Busemeyer et al. 2020; Khrennikov 2023; Asano and Khrennikov 2026). These refinements increase flexibility and thus the burden of showing that a model is not merely descriptive (Section 17) (Broekaert et al. 2017).

5.3 Order Effects: Why Measurement Sequence Matters

If the questions “Is Clinton honest?” and “Is Gore honest?” are represented by non-commuting projectors P_A and P_B , then $p(A \text{ yes, then } B \text{ yes}) = \|P_B P_A \psi\|^2$ differs from $p(B \text{ yes, then } A \text{ yes})$. Wang and Busemeyer showed that any such model, as long as the questions are represented by projectors and the state is pure, obeys a parameter-free constraint, the QQ equality: $[p(A_y B_n) + p(A_n B_y)] - [p(B_y A_n) + p(B_n A_y)] = 0$ (Wang and Busemeyer 2013). Classical Markov models that produce order effects do not generally satisfy it. Wang, Solloway, Shiffrin and Busemeyer tested the equality in 70 national surveys on question order and found it statistically supported, while anchoring-and-adjustment models failed to reproduce the pattern (Wang et al. 2014). Order effects also arise in inference, where evidence presented in different orders leads to different final beliefs (Trueblood and Busemeyer 2011; Trueblood et al. 2017).

Worked Example 1 makes the mechanism concrete (code in Appendix B and Online Resource 2). Take a real two-dimensional belief space with the state at 20° , question A at 0° and question B at 55° . Then $p(A \text{ yes}) = 0.883$ and $p(B \text{ yes}) = 0.671$; $p(A \text{ yes, then } B \text{ yes}) = 0.291$ while $p(B \text{ yes, then } A \text{ yes}) = 0.221$, an order effect of 7 percentage points; and the QQ quantity evaluates to zero to machine precision (Figure 2). The effect disappears exactly when the projectors commute: here the norm of their commutator is 0.66, and setting the two angles equal or orthogonal reduces it to zero.

Two cautions apply. The QQ equality is necessary rather than sufficient evidence for a quantum model: non-Hilbertian and classical models can also satisfy it (Aerts and Sassoli de Bianchi 2017; Kellen et al. 2018), and when interference plays no role the same order effects can be captured by a classical projection model (Moreira and Wichert 2017). The claim therefore rests on the combination of an a priori prediction and its robust confirmation, not on uniqueness.

5.4 The Conjunction Fallacy: Violations of Classical Probability

In the classic Linda problem, people judge “Linda is a feminist bank teller” more likely than “Linda is a bank teller” (Tversky and Kahneman 1983). The quantum account treats the conjunction as a sequence of projections: judging the more representative property first (“feminist”) moves the state towards a subspace from which “bank teller” is more reachable than from the initial state (Busemeyer et al. 2011). Worked Example 2 shows this. With the state at 0° , “feminist” at 40° and “bank teller” at 80° , $p(\text{bank teller}) = 0.030$ but $p(\text{feminist, then bank teller}) = 0.344$, so the model produces a large conjunction fallacy without any free parameter beyond the three angles. The same model predicts disjunction errors and the relative size of both effects (Busemeyer et al. 2011).

The conjunction fallacy is also where the quantum account is most contested. Boyer-Kassem, Duchêne and Guerri derived testable consequences of the main quantum-like models and found experimental results at odds with them (Boyer-Kassem et al. 2016); web-scale tests suggest that emergent meaning in concept combination, rather than order, drives the effect (Aerts et al. 2017); and the “probability theory plus noise” model explains conjunction and related biases with noisy classical reasoning (Costello and Watts 2014, 2018). A review of experimental designs shows that most studies have explored a narrow part of the possible outcome space (Veloz and Sobetska 2024). For autonomous systems, the practical conclusion is that conjunction-fallacy modelling alone is weak justification for adopting QP.

5.5 Interference Effects: Double-Slit Analogy in Decision-Making

In the double-slit experiment, the pattern produced when the path is unobserved is not the sum of the patterns for each path. The decision analogue is the disjunction effect. In a one-shot Prisoner’s Dilemma, Shafir and Tversky found that 97% of participants defected when told the opponent had defected and 84% when told the opponent had cooperated, but only 63% when the opponent’s move was unknown (Shafir and Tversky 1992). The law of total probability requires the unknown case to lie between 0.84 and 0.97 (0.905 with equal priors). In QP the unknown-case probability is the classical mixture plus an interference term $2\sqrt{\pi_1 p_1 \pi_2 p_2} \cos \theta$, where π_1 and π_2 are the prior weights of the two branches (Appendix A.4). Worked Example 3 solves for this term (Figure 4): the observed 0.63 requires an interference term of -0.275 , which corresponds to a phase $\theta \approx 108^\circ$ (Pothos and Busemeyer 2009). The model explains the violation, but in this form the phase is fitted to the very datum it explains, which is why later work has sought ways to predict it, as discussed in Section 7.3 (Moreira and Wichert 2016; Wichert et al. 2020). Yukalov and Sornette’s quantum decision theory provides an alternative formalism in which an “attraction factor” plays the role of interference and a quarter-law gives an a priori estimate of its size (Yukalov and Sornette 2011, 2016).

5.6 Evidence: Why Quantum Models Outperform Classical Approaches

The strongest evidence comes from three kinds of test. The first is a priori, parameter-free predictions confirmed in new data, above all the QQ equality (Wang et al. 2014). The second is dynamic signatures that Markov models cannot produce: interference of an intermediate choice on later confidence (Kvam et al. 2015), oscillations in preference strength (Kvam et al. 2021) and a Zeno-like slowing of belief change under frequent measurement (Yearsley and Pothos 2016). The third is quantitative model comparison, where Hilbert-space multidimensional models compared favourably with Bayesian networks (Busemeyer and Wang 2018) and a quantum sequential sampler compared favourably with the Bayesian sampler (Huang et al. 2025a). Contextuality analyses that control for direct influences have also found genuine contextuality in




Fig. 4 Interference in the Prisoner’s Dilemma disjunction effect. Observed defection rates from Shafir and Tversky (1992) and the classical prediction from the law of total probability with equal priors. The observed rate for the unknown condition lies outside the classically admissible range; QP accommodates it with an interference term of -0.275 (phase $\approx 108^\circ$). Values computed with the code in Online Resource 2

carefully designed tasks (Cervantes and Dzhafarov 2018; Basieva et al. 2019), and a temporal Bell inequality was violated in one of the tested cases (Waddup et al. 2023).

The picture is not one-sided, and the section title should be read with that qualification. When the contextuality-by-default method was applied to earlier datasets, it found no true contextuality once direct influences were accounted for (Dzhafarov et al. 2016). In dynamic belief-change data, quantum advantages were confined to low-coherence conditions and shrank against a stronger Markov baseline (Busemeyer et al. 2019). An open-system model of choice and response time was needed for only a subset of participants, the rest being fitted adequately by a Markov model (Epping et al. 2023). Of the 71 theory and evidence studies in our corpus, 37 test their models on human data (E3) and 31 are formal models without data (Table 1 and Online Resource 1). The fair summary is therefore that QP models outperform standard classical models on a well-defined class of phenomena involving incompatible, sequential or context-dependent judgements, and are unnecessary elsewhere. This class is exactly the one autonomous agents encounter when they interact with people.

6 Tier 1 (Deep) — Perception, Memory and Decision-Making in Autonomous Systems

Sections 6–12 are organised in three tiers of depth: Tier 1 (this section) treats the functions with the most developed models, Tier 2 (Section 7) the functions with partial evidence, and Tier 3 (Section 12) the speculative multi-agent applications. Figure 5 maps the QP principles of Section 5 onto these agent functions. Three questions recur: whether the judgements involved are incompatible (so that a single joint distribution does not exist), whether the agent’s task requires predicting or aligning with such judgements, and whether the QP model makes an a priori prediction that could fail.

Only when all three hold is a QP module likely to add value over a classical model; Section 14 turns this into a decision guide.




Fig. 5 From quantum-probability principles to cognitive phenomena and to the agent functions reviewed in this survey

6.1 Perception Under Sensor and Feature Ambiguity

6.1.1 Problem: Non-Independent Sensor and Feature Measurements

Classical perception pipelines assume that each percept can be assigned a probability within one joint distribution over features. This assumption fails when features are not jointly measurable. Examples include ambiguous figures, vague categories whose boundaries shift with context (Blutner et al. 2013), similarity judgements that are asymmetric and order-dependent (Pothos et al. 2013), and robots whose limited sensors cannot resolve mutually exclusive interpretations. In these cases the order in which features are queried, or the framing of the query, changes the outcome.

6.1.2 Quantum Solution: Entanglement and Superposition in Feature Spaces

Lanza, Solinas and Mastrogiovanni proposed the first explicitly quantum-like robot perception model: a percept is a qubit whose amplitudes encode the balance between two mutually exclusive interpretations, and perception is a measurement (Lanza et al.

2020). The multi-sensory extension encodes n normalised sensor readings in n qubits and defines “query operators” that let the robot read out its degree of belief in a target world state (Lanza et al. 2021). The work is a formal proof of concept evaluated in simulation, and runs on real quantum hardware produced large errors. No comparison with Bayesian or fuzzy fusion was reported, so the models demonstrate representability rather than improved perception. The quantum geometric model of similarity (Pothos et al. 2013) and quantum-like concept representations (Aerts 2009; Aerts et al. 2013) supply the cognitive templates such robot models could follow, and quantum-like Bayesian networks provide an inference engine for them (Section 7.3) (Moreira and Wichert 2014).

6.2 Memory and Episode Recognition

6.2.1 Problem: Context-Dependent Retrieval and Associations

Human memory is reconstructive: whether an item is recognised depends on the retrieval cue and on earlier retrievals, and people can simultaneously endorse mutually exclusive memory claims (episodic over-distribution). A classical memory store with fixed item probabilities cannot produce these patterns without additional mechanisms.

6.2.2 Quantum Solution: Superposed Memory States and Contextual Collapse

The generalised quantum episodic memory model treats verbatim and gist traces as incompatible features and accounts for over-distribution and subadditivity (Trueblood and Hemmer 2017). Busemeyer, Ozawa, Pothos and Tsuchiya compared three ways of incorporating episodic memory into quantum judgement models, which differ in whether memory updates act like belief updates or like non-commuting projections (Busemeyer et al. 2025). Associative structure in the mental lexicon also has a quantum-like description (Bruza et al. 2009), and order effects in the memory and belief queries of intelligent agents have been analysed in the same terms (Kak 2018). For robots, the analogy is to context-dependent episodic retrieval in long-lived agents, for example recognising a place or a person under changing cues. No robot memory system based on these models has been reported. This is an open gap: the value of a QP memory for an artificial agent is unestablished, although the models suggest how an agent could reproduce, and therefore anticipate, human recall behaviour.

6.3 High-Stakes Decision-Making Under Uncertainty

6.3.1 Problem: Violation of the Independence Axiom in Decisions

Expected-utility theory and the POMDP formalism rest on Savage’s sure-thing principle (Savage 1954). Humans violate it (Section 5.5) (Shafir and Tversky 1992), and the Ellsberg paradox shows that ambiguity, not just risk, shapes choice (Ellsberg 1961). When an autonomous system must predict or collaborate with human decisions (in driving, triage, defence or shared control), a normative model will mispredict precisely those choices.

6.3.2 Quantum Solution: Interference Effects in Sequential Decisions

Quantum decision models represent the unresolved state of the world as a superposition, so that deciding with the uncertainty unresolved differs from deciding in each resolved case (Pothos and Busemeyer 2009; Busemeyer and Bruza 2012). Quantum decision theory treats prospects as superpositions and derives a quarter-law for the attraction factor (Yukalov and Sornette 2011, 2016). Its evolutionary extension addresses time-inconsistent choices (Yukalov 2020), and a quantum-conceptual account of expected-utility violations has been proposed (Aerts et al. 2011a). Interference also arises from an intermediate categorisation (Wang and Busemeyer 2016) or choice (Kvam et al. 2015), a pattern directly relevant to agents that must first classify a situation and then act on it. The most direct robotics application is Song et al.’s quantum-like Bayesian model of pedestrian crossing decisions for automated driving: the model captures crossings that a classical model treats as irrational and covers edge events in interactive scenes better than a Social-LSTM baseline (Song et al. 2022). That study is a simulation on trajectory data (E4), and its interference parameters appear to be set rather than learned. It nonetheless marks the point at which quantum cognition entered a safety-critical autonomy problem.

7 Tier 2 (Medium) — Learning, Planning and Sensor Fusion

7.1 Quantum Reinforcement Learning in Autonomous Agents

7.1.1 Exploration vs. Exploitation Trade-off

Every RL agent must balance exploiting what it knows against exploring what it does not. Classical agents use ϵ -greedy, softmax or optimism-based rules (Sutton and Barto 2018). Quantum-inspired RL replaces these with a representation in which each state holds a superposition of actions and action selection is a measurement.

7.1.2 Quantum Superposition for Policy Search and Value Functions

Dong, Chen, Li and Tarn’s quantum RL (QRL) stores, for each state, amplitudes over “eigen-actions”; rewards drive Grover-style amplitude amplification of the chosen action, so the exploration probability adapts automatically (Dong et al. 2008). The robust quantum-inspired variant was validated on a real mobile robot and was more robust to the learning rate than temporal-difference learning (Dong et al. 2012), and quantum-inspired experience replay extends the idea to deep RL (Wei et al. 2022). These algorithms are classical and should be judged against modern classical exploration strategies. Our illustrative experiment in Section 16.3 does exactly this and finds a clear trade-off: faster early learning with noisy rewards, but no stable convergence. Projective simulation is a different, physics-motivated agent model in which deliberation is a random walk through episodic memory (Briegel and De las Cuevas 2012). Its convergence in Markov decision processes has been proved (Boyajian et al. 2020), a

photonic hardware architecture has been proposed for it (Flamini et al. 2020), a photonic experiment demonstrated a quantum speed-up in learning (Saggio et al. 2021), and its classical version has been benchmarked against Q-learning and SARSA (Melnikov et al. 2018) and used on real robots (Section 11.2). Variational-circuit RL, which uses quantum computation rather than QP as a model of cognition, is reviewed in the companion survey (Section 10.2).

7.2 Planning Under Uncertainty

7.2.1 Classical Planners and State Representation Limitations

Classical planners and POMDP solvers represent belief as a distribution over a fixed state space (Lauri et al. 2023). They cannot represent the context-dependent preferences and goals of human partners, which may be constructed during interaction rather than retrieved (Kvam et al. 2021).

7.2.2 Quantum Planning: Superposed Goal States and Interference in Paths

Two different ideas are at work. Quantum algorithms that search plan spaces faster belong to quantum computing and are reviewed in the companion survey (Guru and Chauhan 2026); they inherit the limits of noisy hardware (Preskill 2018). The cognitive idea uses quantum dynamics to model deliberation. Quantum walks over option networks can produce choice probabilities that violate the sure-thing principle (Martínez-Martínez and Sánchez-Burillo 2016), Lindblad models describe how a deliberating agent escapes classical equilibria (Asano and Khrennikov 2026), and a maze example contrasts classical and quantum decision scenarios for an intelligent agent (Kak 2018). Such models could let a planner anticipate how a human collaborator’s goals will shift, but none has yet been integrated into a robot planner. We found no evidence that QP planning improves plan quality for autonomous systems, and we regard this as an open research question rather than an established capability.

7.3 Multi-Source Fusion with Non-Independent Modalities

7.3.1 Problem: Kalman Filter and Classical Fusion Assumptions Break Down

The Kalman filter is optimal for linear-Gaussian systems with known noise (Kalman 1960), and its non-linear and learned variants dominate robotic state estimation (Haykin 2001; Thrun et al. 2005). Its assumptions (a single joint distribution, independent noise, order-invariant updates) are appropriate for physical sensors measuring the same physical quantity. They break down when “sensors” are judgements from different contexts, for example a human report, a language instruction and a vision classifier, whose outputs cannot be embedded in one joint distribution. It is important to be precise here: for physical pose estimation from independent Gaussian sensors, QP adds nothing and the Kalman filter remains optimal (Section 16.1).

7.3.2 Quantum Fusion: Entangled State Updates and Measurement Context

Quantum-like Bayesian networks (QLBNs) replace the conditional probabilities of a Bayesian network with complex amplitudes, so marginalising over an unobserved node produces interference terms (Moreira and Wichert 2014, 2016). The central difficulty is that every unobserved node adds a free phase. Successive work has proposed heuristics, including a similarity heuristic (Moreira and Wichert 2016), a “law of balance” and “law of maximum uncertainty” (Wichert et al. 2020) and influence diagrams that add utilities (Moreira et al. 2020b), to set the phases without fitting them to the data being predicted. Extensions combine QLBNs with belief entropy for multi-attribute decisions (She et al. 2021), with Dempster–Shafer evidence theory (Gao et al. 2025), and with entanglement to capture biased selection, where a biased entangled model achieved an RMSE of 3.4 on Prisoner’s Dilemma data, the best among the models compared (Meghdadi et al. 2022). In multimodal ML, quantum-inspired fusion treats each modality’s decision as a measurement in an incompatible basis; on video sentiment benchmarks this outperformed decision-level and content-level fusion and handled cases in which every unimodal classifier failed (Gkoumas et al. 2021; Li et al. 2021a). These results are on benchmark labels, not human decision data, and the incompatibility assumption itself was not tested.

Part C — Applications and Validation (Sections 8–16)

8 Quantum Cognition for Machine Learning and AI

8.1 Feature Learning and Representation Problem

8.1.1 Classical Approach: Dimensionality Reduction (PCA, t-SNE)

Classical representation learning maps inputs to real-valued feature vectors, whether by linear projections such as PCA, non-linear embeddings such as t-SNE, or learned deep features (Krizhevsky et al. 2012). Probabilities are then obtained by a fixed read-out (for example softmax), which produces a single joint distribution over labels regardless of the order or context of the query.

8.1.2 Quantum Approach: Superposed Feature Spaces and Entangled Representations

Three distinct ideas appear in the literature and should not be conflated. Quantum kernel methods map data into a quantum feature Hilbert space and require quantum hardware for any advantage (Schuld and Killoran 2019; Biamonte et al. 2017). Their advantages are fragile, because trainability limits (Cerezo et al. 2021) and dequantisation results (Tang 2019; Aaronson 2015) often remove the claimed speed-up. Quantum-inspired representations use complex-valued vectors and density matrices on classical hardware: a word is a superposition of semantic “sememes”, a sentence is a mixture, and matching is a measurement (Li et al. 2019; Zhang et al. 2018, 2025). Finally, quantum-cognitive representations aim to reproduce human-like judgement

patterns rather than to improve accuracy (Maksimovic and Maksymov 2025). Tensor-network classifiers, which are classically simulable quantum-style models, reached below 1% test error on MNIST (Stoudenmire and Schwab 2016). Quantum-inspired matching models have typically performed comparably to strong CNN and RNN baselines rather than better, while offering representations their authors describe as more interpretable (Li et al. 2019). Those interpretability claims have not been tested with users.

8.2 Classification Under Label Ambiguity

8.2.1 When Multiple Labels Apply Simultaneously: Order-Dependent Classification

Real categories are vague and overlapping, and human annotators’ labels depend on which category is considered first (Blutner et al. 2013; Wang and Busemeyer 2016). A softmax classifier assigns a single distribution over labels and cannot represent the fact that the answer to “is it a cat?” depends on whether “is it wild?” was asked first. For ML systems trained on human labels, or used alongside human decision-makers, this is a modelling mismatch rather than label noise.

8.2.2 Quantum Cognition: Interference in Multi-Class Decisions

Representing labels as non-commuting projectors on a learned state produces order-dependent label judgements (Case Study 2, Section 16.2). Quantum-cognitively motivated decision fusion used this structure to combine modality-specific sentiment judgements (Gkoumas et al. 2021), and a survey of quantum-cognitively inspired sentiment analysis maps the design space (Liu et al. 2023). Quantum-cognitive neural networks that use tunnelling-inspired activations reproduced human-like confidence and uncertainty on image classification and trained up to 50 times faster than their classical counterpart. That human-likeness was assessed by simulation rather than against participant data (Maksimovic and Maksymov 2025). No study has yet shown that order-aware QP classifiers improve agreement with individual human annotators on a standard benchmark; this is a clear, testable gap (Benchmark 3, Section 22).

8.3 Knowledge Representation and Semantic Similarities

Quantum logic has a long history in semantic representation. Widdows and Peters showed that negation and disjunction on word vectors behave as quantum-logical connectives, allowing unwanted word senses to be removed from queries (Widdows and Peters 2003). Density-matrix language models captured term dependencies for retrieval (Sordoni et al. 2013), and a survey of quantum-theory-inspired information retrieval covers representation, ranking and user cognition (Uprety et al. 2020). Widdows, Kitto and Cohen give the broadest account of quantum mathematics in AI (Widdows et al. 2021). Compositional quantum NLP (DisCoCat) was run on real quantum hardware for small sentence-classification tasks, a feasibility demonstration with small datasets, long queue times and no quantum advantage (Lorenz et al. 2023); two reviews cover the field (Widdows et al. 2024; Nausheen et al. 2025). Contextual

word embeddings have been recast in terms of quantum contextuality, as a framework without an empirical evaluation (Svozil 2025). For agents, such as robots grounding language instructions or dialogue systems tracking context, the relevant contribution is a principled way to represent meaning that depends on context. The contribution does not come from quantum computation.

9 Quantum Cognition for Deep Learning

9.1 Activation and Representation Learning

9.1.1 Classical DL: Neuron Activations as Probabilities

In standard deep networks, activations are real numbers and output probabilities come from a softmax. Uncertainty is typically estimated post hoc (ensembles, dropout, Bayesian layers), and all outputs are assumed to live in one classical probability space.

9.1.2 Quantum DL: Superposed Activations and Non-Classical Interference

Complex-valued, quantum-inspired networks carry phases as well as magnitudes, so that intermediate representations can interfere. CNM encodes words as complex superpositions and sentences as density matrices and matched strong baselines on question answering (Li et al. 2019). A quantum many-body wave-function view of language modelling offered a theoretical justification for CNN structure (Zhang et al. 2018), and a recent sequence model with unitary Hamiltonian dynamics and Born-rule decoding proves an expressivity gain on a constructed task family but has not yet been tested on language data (Nebli et al. 2026). Quantum-inspired networks for conversational emotion recognition (Li et al. 2021b) and multimodal sentiment (Li et al. 2021a) performed on a par with the state of the art. Quantum Bayesian neural networks propose quantum algorithms for posterior sampling but were evaluated only in classical simulation (Berner et al. 2021). The consistent pattern is parity, not superiority, on accuracy, with claims of interpretability that remain to be tested.

9.2 Backpropagation and Gradient Descent Under Uncertainty

9.2.1 Problem: Gradient Vanishing and Loss-Function Plateaus

Deep networks face vanishing gradients; variational quantum circuits face the more severe problem of barren plateaus, in which gradients vanish exponentially with the number of qubits for global cost functions (Cerezo et al. 2021). This is a limitation of quantum hardware models, not of quantum-like models run classically, which are trained with ordinary automatic differentiation (Bergholm et al. 2018; Broughton et al. 2020).

9.2.2 Quantum Gradient Descent: Superposed Weight Updates

We found no evidence that “superposed weight updates” improve the training of classical deep networks, and we caution against the claim. Where quantum circuits are

trained, gradients are estimated by parameter-shift rules and inherit barren-plateau and noise limits (Cerezo et al. 2021; Preskill 2018). Where quantum-inspired networks are trained classically, their complex-valued parameters are optimised by standard gradient descent. The contribution of quantum cognition to optimisation is therefore conceptual, for example modelling how a learner’s uncertainty about incompatible hypotheses evolves, rather than algorithmic.

9.3 Attention Mechanisms and Multi-Head Architectures

9.3.1 Order-Dependent Attention: Measurement Sequence in Transformer Models

Transformer outputs depend on prompt order and framing. Human-like response biases in survey-style prompts were found only partially, and inconsistently, across nine LLMs (Tjauatja et al. 2024), while cognitive-psychology batteries revealed both human-like and non-human-like reasoning in GPT-3 (Binz and Schulz 2023). Quantum-cognition methods give these observations a precise vocabulary and test statistics. Lo, Sadrzadeh and Mansfield constructed a linguistic schema modelled on a contextual quantum scenario and, using a masked language model, found about 0.15% of instances contextual under the sheaf-theoretic criterion and about 71% under contextuality-by-default (Lo et al. 2025). Context effects in LLM similarity judgements mirror human order effects in some settings but not others (Uprety et al. 2024). Bell-type tests on LLM concept combinations reported violations and Bose–Einstein-like word statistics resembling human data (Aerts et al. 2026). GPT-5 committed far fewer conjunction fallacies than humans, and its judgement pattern was closer to early QP models than to the noisy variants that fit people (Imannezhad et al. 2026). An audit of question-order effects with the QQ equality found near-deterministic saturation in a small model, showing that the QQ test can be uninformative when an LLM’s answers are saturated (Kang 2026). Together these studies suggest that attention-based models exhibit some context effects that QP models describe well, but they do not show that transformers are quantum-like in general. The analogy between attention and “soft measurement” remains a hypothesis to be tested (Section 14.3).

10 Quantum Cognition for Reinforcement Learning

10.1 Markov Decision Process (MDP) Assumptions and Their Violations

10.1.1 Problem: Memory Dependencies and Context Effects

RL assumes the Markov property and stationary, additive rewards (Sutton and Barto 2018). Human feedback, which increasingly shapes RL for robots and LLMs, violates both: evaluations depend on what was asked before, preferences oscillate over time (Kvam et al. 2021), and trust evolves non-monotonically (Roeder et al. 2023). Reproducibility studies also show how sensitive deep RL is to seeds and hyper-parameters, which sets a high bar for claims of improvement (Henderson et al. 2018).

10.1.2 Quantum RL: Non-Markovian Memory and State Superposition

Open-system quantum–Markov models provide a principled way to model non-Markovian, context-dependent reward sources (Busemeyer et al. 2020; Epping et al. 2023). To our knowledge no RL algorithm has yet used such a model of the human reward-giver, which we identify as a promising direction (Section 24).

10.2 Policy Search and Exploration Strategies

10.2.1 Classical RL: Epsilon-Greedy, Softmax, Boltzmann Exploration

Classical exploration strategies are simple, well understood and hard to beat on small problems (Sutton and Barto 2018), and deep RL for real robots has succeeded mainly through careful engineering and simulation-to-real transfer rather than new exploration principles (Lauri et al. 2023).

10.2.2 Quantum RL: Quantum Superposition for Simultaneous Action Evaluation

Three lines coexist, and only the first is in the scope of this review. Quantum-inspired RL (Section 7.1) is classical: action amplitudes are “measured” to select an action, and amplitude amplification shifts the policy towards rewarded actions (Dong et al. 2008, 2012; Wei et al. 2022). Quantum-agent theory proves speed-ups in deliberation and learning separations on specially constructed tasks (Dunjko and Briegel 2018). Reinforcement learning with parametrised quantum circuits in place of neural networks uses quantum computation as a function approximator rather than QP as a model of cognition; it is reviewed, with its robot studies, in the companion survey (Guru and Chauhan 2026) and in dedicated surveys (Meyer et al. 2022; Chen 2024). Its one consistent benefit there is parameter efficiency, and no study shows better policies than tuned classical networks of comparable compute. The question for this review is different: whether a QP model of how a learner, or a human teacher, holds and revises incompatible evaluations helps an agent learn. No study has yet tested this.

10.3 Value Estimation and Q-Learning Under Uncertainty

10.3.1 Order Effects in Temporal-Difference Learning

Temporal-difference learning is order-sensitive by construction: the same experiences presented in different orders yield different value estimates. Classical RL treats this as a nuisance that experience replay reduces. Quantum cognition suggests an alternative reading when the learner is human-like: order effects in the evaluation of evidence are systematic and predictable (Trueblood and Busemeyer 2011; Trueblood et al. 2017).

10.3.2 Quantum Q-Values: Interference in Reward Prediction

Interference in reward prediction has not been demonstrated as a benefit in any benchmark we found. In our own illustrative experiment (Case Study 3, Section 16.3),

amplitude-based exploration without a mechanism for reducing amplitudes became unstable under reward noise, which highlights a design gap rather than an advantage.

11 Quantum Cognition for Physical AI and Embodied Systems

11.1 Sensorimotor Integration and Body Schema

11.1.1 Problem: Proprioception, Interoception and Multimodal Ambiguity

Embodied agents must integrate signals whose meaning depends on the body’s state and the task: proprioceptive, interoceptive, tactile and visual signals can conflict, and their interpretation depends on what the agent is trying to do. Predictive-processing accounts cast this as prediction-error minimisation over a generative model (Friston 2010; Clark 2013). Classical cognitive architectures implement it with symbolic and probabilistic modules (Laird et al. 2017; Kotseruba and Tsotsos 2020).

11.1.2 Quantum Embodied Cognition: Entangled Sensor-Motor States

Robot-facing proposals fall into four groups (Online Resource 1, Table S1). The first is quantum-like perception models, already discussed (Lanza et al. 2020, 2021). The second is quantum-circuit models of robot “thinking” and control: a classic-quantum dialogue model of robot creativity (Mannone et al. 2024), a teleo-reactive robot whose action rules are evaluated by quantum kernels and measurement-based circuits, reaching about 0.91 accuracy offline on wall-following data (Chella et al. 2026), and a neuro-cognitive agent architecture using Grover and Deutsch–Jozsa routines in a game environment (Daglarli 2025). The third is quantum models of robot emotion and affect: emotion spaces encoded in qubits (Yan et al. 2019), Plutchik-based emotion circuits for companion robots (Yan et al. 2021c), a review of emotion-space models (Yan et al. 2021b), mixed states of reflection and affect realised as small entangling circuits (Hoorn and Ho 2019), the Q-Coppélia affective decision framework for human–robot interaction (HRI), which its authors present as a theoretical exercise because the circuit needs too many qubits even to simulate (Ho and Hoorn 2022), an early quantum-automata model of emotional robot behaviour (Lukac and Perkowski 2007) and the QUADRI social-robotics programme (De Carolis et al. 2025). The fourth is architectural positions (Aerts et al. 2011b; Moreira et al. 2020a). Of the 17 robot-facing studies, 11 are at E1–E2: circuits and formal models without user studies or classical baselines. The circuit-based models also blur the line between quantum-like modelling and quantum computing, which should be made explicit because only the former is deployable today. Figure 6 shows how a quantum-like module can sit in an embodied agent.




Fig. 6 A hybrid architecture for a quantum-like decision module in an autonomous agent. Classical perception and a context encoder feed a QP state; context-dependent measurements (projectors or instruments) produce graded, order-aware judgements that a classical planner or RL policy acts on, with explanation and escalation to a human. Feedback updates the state as an open system

11.2 Affordance Perception and Action Selection

11.2.1 Classical Approach: Discrete Action Sets and Binary Affordances

Classical robots typically treat affordances (what an object allows the robot to do) as discrete, independently detectable properties, and select actions by maximising expected utility over a fixed action set.

11.2.2 Quantum Approach: Superposed Affordances and Interference in Action Space

The most mature embodied line is not quantum cognition but projective simulation (PS) run classically. A real robot learned hierarchical manipulation skills through play, reaching 90% success in about 80 roll-outs in simulation (Hangl et al. 2016), and PS with active learning and behaviour composition solved previously unsolvable tasks and improved learning speed by about 30% (Hangl et al. 2020). Superposed affordances in the QP sense, where an object’s affordances are incompatible and the order of evaluation matters, have not been implemented on a robot. They remain a hypothesis (Section 12.4).

11.3 Predictive Processing and Action Prediction Errors

11.3.1 Order-Dependent Prediction Errors in Physical Interaction

Predictive processing treats perception and action as the minimisation of prediction error (Friston 2010; Clark 2013). Quantum-like models add one ingredient: the order in which predictions are tested can change the result. In physical interaction, such as handover, shared manipulation or co-driving, the human partner’s intentions are themselves constructed during the interaction (Kvam et al. 2021), so a robot that queries (by gesture, speech or motion) changes what it is trying to predict. We know of no experiment that tests this in physical HRI; the handover setting often cited for it has no published quantum-like study, and the related works listed in some bibliographies could not be verified (Section 1.5.2).

12 Tier 3 (Shallow) — Speculation on Multi-Agent Coordination and Collective Cognition

12.1 Quantum Cognition in Multi-Robot Coordination

Quantum game theory showed that entangled strategies can escape classical dilemmas (Eisert et al. 1999). Physical entanglement between robots and quantum communication links belong to the companion survey (Guru and Chauhan 2026); quantum-like approaches instead model interdependent agents. Lawless argues that interdependence in autonomous human–machine teams is quantum-like: team performance is not decomposable into member contributions (Lawless 2020; Lawless et al. 2023; Lawless and Moskowitz 2025). Human interaction data show context sensitivity that a quantum model fitted better than a matched classical model (Waddup et al. 2021). Quantum game models, simulated classically, have been proposed for interaction-aware automated driving with two players (Essalmi et al. 2025) and with several players and strategies, where they raised success rates and reduced collisions against classical game-theoretic models in simulation (Essalmi et al. 2026); an independent reproduction of the earlier model could not confirm its collision results. Multi-robot emotion fusion through entangled qubit encodings remains a formal model (Yan et al. 2021a).

12.2 Distributed Quantum-Like Decisions Across Swarms

Networks of quantum decision-makers exchanging information converge to consensus and reproduce a dynamic disjunction effect (Yukalov et al. 2018, 2022), and affective AI has been framed as a “quantum operation” combining cognition and emotion (Yukalov 2023). Quantum stochastic walks give a network model of decision (Martínez-Martínez and Sánchez-Burillo 2016). PS swarms reproduce collective motion (Ried et al. 2019), foraging-dependent alignment (López-Incera et al. 2020) and queue formation with polarisation near 0.99 (Kong et al. 2026), all in simulation with classical PS.

12.3 Concept Learning and Semantic Emergence

Entangled concept combinations ([Aerts and Sozzo 2011](#)), Fock-space concept dynamics ([Aerts et al. 2013](#)) and Bell-type structure in LLM language ([Aerts et al. 2026](#)) suggest that shared meaning in groups of agents could be modelled with quantum structure. Whether such models help artificial agents develop shared concepts is untested.

12.4 Open Questions for Future Research

- Can a quantum-like model of human teammates improve a robot team’s coordination with people, compared with Bayesian theory-of-mind models?
- Do quantum-like consensus models predict human–robot group decisions better than classical opinion dynamics?
- Are “superposed affordances” measurable in human behaviour during joint action, and do they help a robot anticipate its partner?
- Which multi-agent claims survive when compared with classical models of equal parameter count, as the quantum game models of driving still need to be ([Essalmi et al. 2025, 2026](#))?

13 Prior Work and Positioning

13.1 Quantum Cognition in Psychology

The psychological literature is mature, with a textbook ([Busemeyer and Bruza 2012](#)), a BBS target article ([Pothos and Busemeyer 2013](#)), tutorials ([Yearsley and Busemeyer 2016; Yearsley 2017](#)), reviews ([Bruza et al. 2015; Pothos and Busemeyer 2022; Huang et al. 2025b](#)) and an active critical literature (Sections 5.6 and 14). Its models are fitted to human data, which places 37 of the 71 theory and evidence studies at level E3.

13.2 Quantum Probability in Decision Science and Cognitive Science

Quantum decision theory ([Yukalov and Sornette 2011, 2016](#)), contextual probability ([Khrennikov 2010, 2024](#)), contextuality-by-default ([Dzhafarov and Kujala 2016](#)) and quantum-like heuristics ([Huang et al. 2025c](#)) provide alternative formalisms. A 2015 survey organised quantum-like decision models by formalism ([Ashtiani and Azgomi 2015](#)).

13.3 Application to Robotics, AI, ML, DL and RL (2020–2026)

Online Resource 1 (Table S1) lists every application study with its domain, model type, evidence level, task and main limitation, and Figure 7 summarises the distribution. Of the 67 application studies, 31 are quantum-like, 24 quantum-inspired, 10 quantum-circuit designs or simulations and 2 quantum circuits run on a QPU. Twenty-four are proposals or formal models (E1–E2), 9 are tested on human data (E3), 29 on simulations, ML benchmarks or LLM probes (E4), and 5 reach a physical system (E5). Two

of the five run small cognitive or language models on quantum processors (Widdows et al. 2023; Lorenz et al. 2023), and three are robot studies that use quantum-inspired RL or classical projective simulation (Dong et al. 2012; Hangl et al. 2016, 2020). None of the five tests a quantum-cognition decision model on a robot, and no study reaches E6.




Fig. 7 Evidence levels of the 138 graded studies in the corpus: the 71 theory and evidence studies (top bar) and the 67 application studies by domain (reviews and background works are not graded). E1 conceptual; E2 formal model; E3 human data; E4 simulation or benchmark; E5 physical system; E6 field deployment. No study reaches E6

Human-AI trust and oversight

Human trust in automation is the application domain in which quantum-like models have the strongest human-data support. Quantum models outperformed Markov models in predicting how people’s reliability judgements of an automated system evolve over an interaction (Roeder et al. 2023), and a quantum-inspired model was fitted to air-traffic controllers’ trust in automation, although with a very small sample (Pushparaj et al. 2021). Quantum-like trust models have been proposed for AI-supported decisions, including military decision aids (Humr et al. 2023, 2025; Humr and Canan 2026), and quantum random walks with interaction-dependent Hamiltonians have been fitted to the reliability ratings of 34 participants who worked with an AI image classifier (van der Meer et al. 2025); computational trust between artificial agents has also been given a quantum-like formulation (Ashtiani and Abdollahi Azgomi 2019). These models share a common structure: the person’s trust is a state that is changed by each interaction, so the order in which the agent presents advice, explanations and errors becomes a design variable. For autonomous systems, the practical uses are to anticipate when a person’s reliance will shift, to choose the order and framing of

explanations, and to trigger escalation to human oversight before trust and system capability diverge (Section 19.2). None of these models has yet been evaluated in closed loop with a robot.

Quality of the application studies

The appraisal (Online Resource 1, Table S5) shows where the evidence is weakest. Model specification is the strongest criterion: 47 of the 61 appraised studies state their state space, operators and parameters in full. A classical model was evaluated on the same task in only 24 of the 61, partly in 16, and not at all in 21. Twenty-seven of the 61 give no a priori constraint on their free parameters, typically the interference phases, and validate them on no held-out data; 18 do. Empirical grounding is split: 19 studies use human data or a realistic agent or AI task, 24 use a narrow benchmark or aggregate data from a few classic paradigms, and 18 use toy examples or no data. Of the 14 studies whose code and data statements we could check, 5 released code or data in full or in part. Seven studies discuss their limitations and classical alternatives in depth; most do so briefly.

13.4 Unique Contribution of This Survey

To our knowledge, this review is the first to make the following six contributions.

1. It is the first systematic review, reported to PRISMA 2020, of quantum-cognition models applied to autonomous agents, with every judgement published as data (Section 1.5 and Online Resources 1, 3 and 4).
2. It connects the psychological evidence base to applications in robotics, reinforcement learning, machine learning and deep learning, large language models and human–AI trust within one framework (Sections 5–12).
3. It separates quantum-like, quantum-inspired and quantum-circuit work consistently throughout.
4. It grades every application study on a six-level evidence scale, from conceptual proposals to field deployment, and appraises each against seven quality criteria.
5. It confronts the applications with the critical literature and with classical competitors (Sections 5.6 and 14).
6. It provides case studies, an illustrative experiment, a decision guide, a validation roadmap, a benchmark suite and a replication checklist, with reproducible code (Sections 16, 15 and 22, Appendix B and Online Resource 2).

14 Head-to-Head Comparison — Quantum vs. Classical (Honest Trade-offs)

Table 4 summarises the comparison of quantum-like and classical approaches.

Table 4 Comparison of classical and quantum-like or quantum-inspired approaches

Criterion	Classical (Bayes, Markov, DRL)	Quantum-like or quantum-inspired	Current evidence
Order and context effects	Require extra mechanisms (anchoring, memory)	Arise from non-commuting projectors; a priori QQ constraint	Strong for QQ (Wang et al. 2014); classical accounts possible (Kellen et al. 2018)
Total-probability violations	Cannot be produced without changing the model	Interference terms	Replicated in two-stage tasks (Wang and Busemeyer 2016; Kvam et al. 2015)
Conjunction fallacy	Noise models (probability theory plus noise)	Sequential projection	Contested (Boyer-Kassem et al. 2016; Costello and Watts 2018)
Dynamics of belief and trust	Markov random walks	Quantum walks, open systems	Mixed: quantum better at low coherence (Busemeyer et al. 2019); trust (Roeder et al. 2023)
Physical state estimation	Kalman and particle filters (optimal)	No advantage expected	No comparison published
Classification accuracy	Strong deep-learning baselines	Parity on benchmarks	(Li et al. 2019, 2021b)
RL performance	Tuned deep RL, scalable	QiRL robust to learning rate; unstable under reward noise	No better policies (Dong et al. 2012); Section 16.3
Parameters and flexibility	Well understood	Phases and instruments add flexibility	Risk of over-fitting (Broekaert et al. 2017)
Interpretability	Varies	Geometric parameters; untested with users	No user studies
Hardware needs	CPU or GPU	CPU or GPU (quantum-like); QPU optional	Few-qubit demonstrations (Widdows et al. 2023)

14.1 Quantum Cognition vs. Bayesian Networks

14.1.1 When Classical Wins: Simple, Well-Defined Domains

When all variables are compatible (jointly measurable, context-free), QP reduces to classical probability and adds only parameters (Bruza et al. 2015). For physical state estimation, diagnosis from calibrated sensors (Blanke et al. 2016) and planning with known dynamics, Bayesian networks and filters are simpler, better understood and supported by mature tooling. Markov models remain adequate for most participants in several dynamic tasks (Busemeyer et al. 2019; Epping et al. 2023).

14.1.2 When Quantum Wins: Non-Commutative Measurements and Interference

QP models win where their a priori structure matches the phenomenon: question-order effects that obey the QQ equality (Wang et al. 2014), total-probability violations in two-stage decisions (Pothos and Busemeyer 2009; Wang and Busemeyer 2016) and interference of choice on confidence (Kvam et al. 2015). In direct comparisons, Hilbert-space multidimensional models compared favourably with Bayesian networks (Busemeyer and Wang 2018) and quantum-like influence diagrams predicted

paradoxical choices that expected utility cannot (Moreira et al. 2020b). The main classical competitors are noisy classical reasoning (Costello and Watts 2014, 2018), classical accounts of mirrored order effects (Kellen et al. 2018) and contextuality analyses that attribute effects to direct influences (Dzhafarov et al. 2016). Replies and new tests continue (Busemeyer and Wang 2017; Basieva et al. 2019).

14.2 Quantum Cognition vs. Deep Reinforcement Learning

14.2.1 Classical DRL: Scalable but Data-Hungry

Deep RL scales to high-dimensional robot control but requires large amounts of data and careful tuning, and its reported gains are fragile (Henderson et al. 2018).

14.2.2 Quantum RL: Efficient but Interpretability Challenge

Quantum-inspired RL is cheap and can be robust to learning-rate choices (Dong et al. 2012), but our illustrative experiment shows instability under reward noise (Section 16.3). Circuit-based RL, reviewed in the companion survey, is parameter-efficient but not yet better or faster in wall-clock terms (Guru and Chauhan 2026). The claim that quantum RL is more interpretable is not supported by any user study we found.

14.3 Quantum Cognition vs. Transformer Attention Models

14.3.1 Attention as Soft Measurement: Order Effects in Transformers

Attention weights a set of tokens by their compatibility with a query, loosely analogous to a context-dependent measurement. The analogy is suggestive, and QP tools (the QQ equality, contextuality-by-default) are useful auditing instruments for LLMs (Lo et al. 2025; Kang 2026). However, transformers are classical functions of their input; order effects in LLMs can arise from training data and position encodings without any QP structure, and their human-likeness varies across models and prompts (Tjuatja et al. 2024; Imannezhad et al. 2026). QP models are best used here to describe and test LLM behaviour, not as a claim about its mechanism.

14.4 Hybrid Approach: When to Use Each

Figure 8 condenses the comparison into a decision guide. Keep a classical core for perception, estimation and control, where classical methods are optimal or nearly so. Add a quantum-like module only where the agent must predict, explain or align with judgements that are sequential, ambiguous or context-dependent. Typical cases are human intentions, trust and preferences, and the answers of other learned agents. Validate the module against a classical model with the same number of parameters and against a noise-augmented classical model.




Fig. 8 Decision guide for choosing between classical and quantum-like models in an autonomous system

15 Experimental Validation Roadmap

Because the literature contains almost no controlled comparisons on autonomous systems, we propose a two-phase validation programme. Each task has a classical baseline, a quantum-like model and an a priori QP prediction that can fail.

15.1 Phase 1 (Simulation): Benchmark Tasks

15.1.1 Task 1: Fault Diagnosis in Simulated Autonomous Systems

A simulated mobile robot with injected faults receives diagnostic evidence from heterogeneous sources (sensor residuals, operator reports, log classifiers) presented in different orders. The baseline is a Bayesian-network or observer-based diagnoser (Blanke et al. 2016). The test is whether a QLBN predicts operators’ order-dependent diagnoses better, and whether exploiting it improves joint human–robot diagnosis accuracy. The quantum fault-tolerant filter (Gao et al. 2021) is not an appropriate baseline because it addresses quantum systems.

15.1.2 Task 2: Multi-Sensor Fusion Under Ambiguity

Fusion of physical sensors (baseline: Kalman or particle filter) combined with semantic cues from language or human reports. The prediction is that QP fusion helps only in the semantic, context-dependent channel; it should do no better than the Kalman filter on physical channels (Section 16.1).

15.1.3 Task 3: Navigation and Decision-Making with Conflicting Information

Navigation where route advice from a human, a map and a learned model conflict. Baselines are POMDP planners and ensembles; the quantum-like model predicts the human’s acceptance of advice as a function of order and framing (Song et al. 2022; Roeder et al. 2023).

15.1.4 Task 4: ML/DL Classification Under Label Ambiguity

Classification of ambiguous images, with human labels collected under controlled question orders. The model predicts order-dependent label distributions and the QQ statistic; the baseline is a calibrated softmax classifier with annotator noise.

15.2 Phase 2 (Real Systems): Hardware Validation

15.2.1 Robotic Platforms: Grasping, Manipulation, Navigation

Human–robot handover, shared manipulation and guided navigation are natural platforms because a human judgement is embedded in the loop. The primary outcome is task success and human trust, not model fit alone.

15.2.2 Real-World AI/ML Systems: Object Recognition, Intent Prediction

Decision-support systems in which people accept or reject AI recommendations, as in air-traffic control (Pushparaj et al. 2021) or military decision aids (Hunr and Canan 2026), offer realistic settings with measurable trust dynamics. Pedestrian-intent prediction for automated driving (Song et al. 2022) and object recognition with human-labelled ambiguous categories are the corresponding perception tasks.

15.3 Metrics and Baselines

Table 5 lists the metrics. Every study should report predictive accuracy on held-out data (not only fit), model complexity (parameters, AIC/BIC or cross-validated likelihood), a parameter-matched classical baseline, a noise-augmented classical baseline such as probability theory plus noise (Costello and Watts 2014), and at least one a priori QP test (the QQ equality, a total-probability test or a contextuality test with direct influences controlled) (Dzhafarov and Kujala 2016). The appraisal of Section 13.3 shows that most published applications fall short on exactly these points.

16 Case Studies and Worked Examples

The three case studies use the code in Appendix B (Online Resource 2); their numerical values are collected in Appendix C.

Table 5 Metrics and baselines for validating quantum-like models in autonomous systems

Metric	What it measures	Why it is needed
Held-out predictive log-likelihood AIC, BIC and parameter count	Out-of-sample prediction of human or agent judgements Model complexity	Avoids rewarding flexible models that only fit QP phases and instruments add freedom
QQ statistic or total-probability test	An a priori QP prediction	Makes the QP model falsifiable
Contextuality-by-default with direct influences	Genuine contextuality	Separates contextuality from signalling
Task success, return or error	Autonomous-system performance	The outcome that matters for deployment
Trust and acceptance	Human reliance on the agent	Main target of human-facing QP models
Seeds and confidence intervals	Statistical robustness	RL results are seed-sensitive
Latency and compute	Real-time feasibility	Embedded deployment

16.1 Case Study 1: Pose Estimation from Noisy Multimodal Sensors

16.1.1 Problem Setup

A robot estimates its position along a corridor from two sensors, a lidar reading $z_1 = 2.10$ m (variance 0.04) and a camera estimate $z_2 = 1.90$ m (variance 0.09). A human operator also reports whether the robot has “passed the doorway”, and a vision classifier gives its own judgement.

16.1.2 Classical Approach (Kalman Filter)

The Kalman update gives $\hat{x} = 2.038$ m with variance 0.028 m^2 , and the result is identical whichever sensor is processed first (Kalman 1960). For this physical sub-problem, the Kalman filter is optimal and a QP model cannot improve on it.

16.1.3 Quantum Cognition Approach

The QP model is applied only to the semantic judgement “has the robot passed the doorway?”, which combines two context-dependent cues: the operator’s report ($p = 0.80$ of “passed” when that cue is attended) and the classifier’s judgement ($p = 0.30$). With equal weights the classical mixture is 0.55. A QLBN adds an interference term $2\sqrt{0.5 \cdot 0.80 \cdot 0.5 \cdot 0.30} \cos \theta = 0.49 \cos \theta$, giving 0.55 at $\theta = 90^\circ$ (the classical case), 0.795 at $\theta = 60^\circ$ and 0.305 at $\theta = 120^\circ$; at $\theta = 0$ the unnormalised value exceeds 1, which is why QLBNs must renormalise (Moreira and Wichert 2016; Wichert et al. 2020). The phase encodes how the two cues interact in the human’s judgement, and it must be estimated from, or better predicted for, human data.

16.1.4 Results and Comparison

The worked example shows the correct division of labour: physical fusion stays classical, and QP models the human-facing judgement. We found no published study that

compares QP and classical fusion on robot pose estimation, and we do not expect QP to help there.

16.2 Case Study 2: Image Classification Under Label Ambiguity (ML/DL)

16.2.1 Problem Setup: Objects with Multiple Valid Labels

An image of a large domestic-looking cat is judged by annotators on two questions: “cat?” and “wild animal?”. Human answers to such overlapping categories are order-dependent (Wang and Busemeyer 2016).

16.2.2 Classical Approach (CNN + Softmax)

A CNN with independent sigmoid outputs (Krizhevsky et al. 2012) predicts $p(\text{cat})$ and $p(\text{wild})$ and, if a joint is needed, assumes conditional independence. The joint probability is the same whichever question is asked first.

16.2.3 Quantum Cognition Approach

Represent the image as a state at 35° in a two-dimensional judgement space, with “cat” at 0° and “wild” at 60° . Then $p(\text{cat}) = 0.671$ and $p(\text{wild}) = 0.821$; $p(\text{cat, then wild}) = 0.168$ but $p(\text{wild, then cat}) = 0.205$. The model predicts a specific order effect and, through the QQ equality, a constraint on the four sequential answer patterns that can be tested on annotation data.

16.2.4 Results and Accuracy Gains

No accuracy gains are claimed. The value of the QP model is that it predicts the order-dependence of human labels, which a softmax classifier cannot represent; whether modelling this improves learning from human labels is untested and is the purpose of Benchmark 3 (Section 22).

16.3 Case Study 3: RL Policy Learning Under Reward Uncertainty

16.3.1 Problem Setup: Ambiguous Reward Signals

A tabular agent must reach the goal of a 6×6 grid with four internal walls (shortest path: 10 steps). Reaching the goal gives reward +1 and each step -0.01 . In the noisy condition, Gaussian noise ($\sigma = 0.5$) is added to every reward. Each agent is trained for 300 episodes (maximum 200 steps each) with 30 random seeds.

16.3.2 Classical Approach (Q-Learning)

Because the state space is small, we use tabular Q-learning (learning rate 0.1, discount 0.95) with either ε -greedy exploration (ε decaying from 0.3 to 0.01) or softmax exploration (temperature decaying from 0.2 to 0.01). Deep Q-learning would behave similarly on this task at much higher cost.

16.3.3 Quantum RL Approach

The quantum-inspired agent follows Dong et al.’s QRL (Dong et al. 2008): each state holds four action amplitudes (initially equal), actions are chosen by “measurement” with probability $|\text{amplitude}|^2$, state values are learned by TD(0), and the chosen action’s amplitude is amplified by $L = \lfloor k(r + V(s')) \rfloor$ Grover iterations, capped at three, with $k = 30$ selected from $\{10, 30, 100\}$.

16.3.4 Results and Convergence Analysis

With deterministic rewards, all three agents learned the optimal 10-step path in every seed (Figure 9). The ε -greedy agent converged fastest (first episode at which the 10-episode moving average fell to 15 steps or fewer: 60 ± 6), while QRL (134 ± 20) and softmax (128 ± 10) were similar. With noisy rewards, QRL learned fastest at first: its median path length was 70 steps at episode 100, against 153 for ε -greedy and 200 (no progress) for softmax. However, it did not stabilise. Over the last 50 episodes its mean path length was 67 ± 40 steps, against 14 ± 13 for ε -greedy and 16 ± 17 for softmax, and 5 of 30 seeds never met the convergence criterion. Varying k did not remove the instability (final means of 78, 67 and 74 steps for $k = 10, 30$ and 100). The reason is structural: amplitude amplification only increases the probability of rewarded actions, so an action amplified by a noisy positive reward is never de-amplified. The experiment is small and illustrative, but it shows why quantum-inspired RL must be compared with tuned classical baselines under realistic noise, and it points to a concrete design fix: amplitude updates that can also decrease (Appendix E). Full results are in Appendix C.

Part D — Grounding and Implications (Sections 17–20)

17 Limitations and Failure Modes

17.1 When Quantum Cognition Does Not Help

QP models do not help when all relevant variables are compatible, which covers most physical sensing, estimation and control; the models then reduce to classical probability with extra parameters (Bruza et al. 2015). They do not help when the target effects are weak or absent in the population of interest: several conjunction-fallacy predictions failed (Boyer-Kassem et al. 2016), classical noise models explain many judgement biases (Costello and Watts 2014, 2018), and in dynamic tasks many participants are fitted as well by Markov models (Epping et al. 2023). Nor do they help when the goal is simply to maximise an agent’s task reward, where no study has shown a QP decision module outperforming a tuned classical one. Finally, they do not help when the effects are artefacts of the measurement interface, for example saturated LLM outputs that render the QQ test uninformative (Kang 2026).




Fig. 9 Case Study 3: steps to reach the goal in a 6×6 grid for ϵ -greedy Q-learning, softmax Q-learning and quantum-inspired RL (Dong et al. 2008), with deterministic (left) and noisy (right) rewards. Lines show the median of 30 seeds of the 10-episode moving average; shading shows the interquartile range; the dotted line is the optimal path length. Experiment run for this review (code in Online Resource 2)

17.2 Interpretability Challenge: “Black Box” Superpositions

QP models are often described as interpretable because their parameters are geometric (angles between subspaces, phases). In practice, complex phases and high-dimensional density matrices are no easier for users to understand than neural weights, and interpretability claims for quantum-inspired networks have not been tested with users (Li et al. 2019, 2021b). Interpretable quantum ML is an active research goal rather than an achievement (Flamini et al. 2024).

17.3 Computational Overhead vs. Accuracy Gain

Quantum-like models are cheap at the scale of psychological experiments but scale exponentially if every binary variable becomes a qubit (Section 20). Quantum-hardware versions add queueing, noise and barren plateaus (Preskill 2018; Cerezo et al. 2021), and dequantisation shows that some claimed quantum speed-ups vanish against clever classical algorithms (Tang 2019; Aaronson 2015). Where accuracy gains have been measured, they are typically parity or small improvements on benchmarks (Sections 8–10).

17.4 Real-Time Constraints in Robotics and Embedded AI

A two- to eight-dimensional QP model can be evaluated in microseconds on an embedded CPU, so quantum-like decision modules are compatible with real-time control. Circuit-based implementations on cloud quantum processors are not: queueing and

communication latencies are orders of magnitude larger than control cycles. Real-time use therefore implies classical simulation of the QP model.

18 Assumptions and Caveats

18.1 Core Assumptions: Linearity, Finite Dimensionality, Closure

QP models assume that belief states form a linear space, that a finite (usually small) dimension suffices, and that measurement outcomes are closed under the chosen projectors or instruments. Dynamic models further assume a specific Hamiltonian or Lindblad generator (Busemeyer et al. 2020; Asano and Khrennikov 2026).

18.2 Where These Assumptions May Break Down

Response replicability shows that simple projective models are too restrictive (Khrennikov et al. 2014); quantum instruments fix this at the cost of extra flexibility (Ozawa and Khrennikov 2020, 2021). Some order-effect data appear to be non-Hilbertian altogether (Aerts and Sassoli de Bianchi 2017). Signalling (direct influence of context on marginal probabilities) violates the assumptions under which Bell-type and contextuality tests are interpreted (Dzhafarov and Kujala 2016). Individual differences are often averaged away in aggregate fits (Epping et al. 2023).

18.3 Uncertainty and Knowledge Gaps

Key unknowns for autonomous systems are: whether QP effects in laboratory tasks generalise to operational settings; how to set phases a priori rather than fit them (Wichert et al. 2020); how QP models should be learned from data at scale; and whether modelling human QP effects improves joint human-agent performance. A further source of uncertainty for this review is that 51 of the 67 application studies were extracted from their abstracts, and several recent preprints have not yet been peer-reviewed (Section 1.5.6).

19 Ethical, Safety and Transparency Implications

19.1 Transparency: Can We Explain Quantum Decisions in AI/ML?

A QP decision module can be explained in terms of its elements: which questions were asked, in which order, and how each answer changed the state. This is a potential transparency advantage over opaque neural policies, but it must be demonstrated with users. Designers should document the ordering of queries, because under QP the ordering is part of the decision.

19.2 Safety: Escalation to Human Oversight

Quantum models of trust in AI provide the most direct safety application. They predict when human reliance will shift, which can trigger escalation or explanation (Roeder et al. 2023; Humr et al. 2025; Humr and Canan 2026; van der Meer et al. 2025). Formal verification of autonomous robots remains classical (Luckcuck et al. 2019; Dreossi et al. 2019); a QP module inside a verified system should be bounded by classical safety envelopes, so that its context-dependence cannot drive the system outside certified behaviour.

19.3 Accountability: Who is Responsible When Autonomous Systems Fail?

Because QP models make outcomes depend on the sequence of interactions, logs must record the order and framing of queries to allow audit. An order-dependent model that is not logged is effectively unaccountable, and responsibility for its failures then cannot be traced to a design decision.

19.4 Bias and Fairness: Does Quantum Cognition Mitigate or Amplify Bias?

QP models reproduce human biases; that is their purpose in psychology. Used to predict people, they can help a system anticipate and correct for biased human inputs. Used to design an agent’s own decisions, they risk importing those biases deliberately. The choice should be explicit: model human bias, do not emulate it, unless emulation is required for natural interaction and is disclosed. Tests of LLMs show that AI systems may display human-like biases, fewer biases or different biases depending on the model (Binz and Schulz 2023; Tjauatja et al. 2024; Imannezhad et al. 2026).

20 Hardware and Implementation Requirements

20.1 Classical Simulation Requirements

20.1.1 Memory Footprint (Exponential in Qubits)

A pure state over n binary variables needs 2^n complex amplitudes (16 bytes each in double precision): 16 KB for $n = 10$, 16 MB for $n = 20$ and 16 GB for $n = 30$. Density matrices square these figures. Practical cognitive models avoid this by using low-dimensional spaces per question set (Busemeyer and Wang 2018) or tensor-network factorisations (Stoudenmire and Schwab 2016).

20.1.2 Computational Complexity and Scaling

Applying a projector or a unitary costs $O(d^2)$ for a dense d -dimensional operator, and a Lindblad step costs $O(d^3)$. For the dimensions used in the literature ($d \leq 16$), these costs are negligible. QLBNs add a phase per unobserved node, so exact inference grows with the number of interference terms, as in classical exact inference (Moreira and Wichert 2014).

20.2 Real Quantum Hardware (Future)

20.2.1 Error Rates and Coherence Times

Current noisy intermediate-scale devices ([Preskill 2018](#)) suffice for the few-qubit circuits of cognitive models ([Widdows et al. 2023](#)), but noise distorts sensitive order-effect circuits and adds latency. There is no computational need for hardware at these sizes; the companion survey quantifies the latency gap for robot control ([Guru and Chauhan 2026](#)).

20.2.2 Qubit Connectivity and Topology

Connectivity limits matter only for large entangled models, which cognitive models do not yet use. Quantum hardware becomes relevant only if QP models are scaled to many interacting variables or used as quantum-agent learners with proven speed-ups ([Paparo et al. 2014](#); [Saggio et al. 2021](#)).

20.3 Practical Implementation: Software Stack and Frameworks (PyQuante, Qiskit, PennyLane)

For quantum-like models, NumPy or SciPy suffices ([Appendix B](#)), with QuTiP for open-system dynamics. For circuit implementations, Qiskit, PennyLane ([Bergholm et al. 2018](#)) and TensorFlow Quantum ([Broughton et al. 2020](#)) integrate with PyTorch and TensorFlow. PyQuante, named in some roadmaps, is a quantum-chemistry package and is not suited to cognitive modelling; we list it here only to prevent that mistake.

Part E — Integration and Future (Sections 21–25)

21 Connection to Neuroscience — Is Quantum Cognition Biologically Grounded?

21.1 Do Brains Actually Use Quantum Mechanics?

The Orch-OR theory proposes quantum computation in neuronal microtubules ([Hameroff and Penrose 2014](#)), and Fisher proposed nuclear-spin qubits in Posner molecules as a possible quantum memory ([Fisher 2015](#)). Tegmark estimated that decoherence times in the brain are far shorter than neural processing times, making sustained quantum computation implausible ([Tegmark 2000](#)). These hypotheses are contested and are not required by quantum cognition.

21.2 Predictive Processing as a Bridge

Predictive processing and the free-energy principle cast cognition as prediction-error minimisation ([Friston 2010](#); [Clark 2013](#)). They provide a classical framework in which the constructive, context-dependent judgements that QP models describe could arise from hierarchical inference in which each query changes the prior for the next.

21.3 Why Quantum Formalism Fits Behavioural Data (Without Requiring Quantum Brains)

Several proposals show how classical neural systems could implement QP operations. Neural oscillators can realise unitary evolution, projection and state reduction (Busemeyer et al. 2017); uncertainty in neuronal firing can be represented as quantum-like superposition (Khrennikov et al. 2018); and generalised probability theory links neuronal networks and oscillations to quantum-like cognition (Khrennikov et al. 2025b,a; Khrennikov and Yamada 2025). These are existence proofs rather than empirical findings, but they remove the need for a quantum brain.

21.4 Implications for AI/ML and Robot Cognition Design

If QP structure arises from classical mechanisms such as context-dependent sampling, oscillatory coding and constructive inference, artificial agents can reproduce it with classical computation. The design question is therefore not whether to use quantum hardware but whether an agent should represent incompatible perspectives explicitly, as QP models do.

22 Standardised Benchmarks and Reproducibility

22.1 Proposed Benchmark Suite (Robotics, ML, RL, Physical AI)

Table 6 specifies four benchmarks that operationalise the validation roadmap (Section 15). Each pairs a classical baseline with an a priori QP test, so that a benchmark can falsify the QP model as well as support it. The full protocol is in Appendix D.

Table 6 Proposed benchmark suite

Benchmark	Task	Classical baseline	QP model and a priori test	Primary metrics
B1 Fault diagnosis	Simulated robot faults; operator judgements under order-randomised evidence	Bayesian network or observer-based diagnosis	QLBN; QQ test on paired questions	Diagnosis accuracy; held-out likelihood
B2 Sensor fusion	Physical sensors plus semantic cues	Kalman or particle filter	QP only on the semantic channel; interference test	RMSE (physical); acceptance prediction
B3 Label ambiguity	Ambiguous images, counterbalanced question order	Calibrated softmax plus annotator model	Projector model; QQ test	Order-effect prediction; calibration
B4 RL under reward uncertainty	Tabular and continuous control with noisy or human rewards	ϵ -greedy, softmax, SAC or PPO	Quantum-inspired RL; parameter-matched comparison	Return; stability over ≥ 30 seeds

22.1.1 Benchmark 1: Fault Diagnosis (Synthetic and Real)

Simulated robot faults with order-randomised evidence presented to human operators, later repeated on a real platform; metrics are diagnosis accuracy, the predictive log-likelihood of operator judgements and a QQ test on paired questions (Task 1, Section 15).

22.1.2 Benchmark 2: Multi-Modal Sensor Fusion

Physical sensors plus semantic cues (Case Study 1); metrics are the estimation error on physical channels, where parity with the Kalman filter is expected, and the prediction of human acceptance on semantic channels.

22.1.3 Benchmark 3: ML Classification Under Label Ambiguity

Ambiguous images labelled under counterbalanced question orders (Case Study 2); metrics are the prediction of order effects, the QQ statistic and calibration against a softmax classifier with an annotator model.

22.1.4 Benchmark 4: RL Policy Learning Under Reward Uncertainty

Tabular and continuous-control tasks with noisy or human-generated rewards (Case Study 3); metrics are sample efficiency, final return, stability across 30 or more seeds (Henderson et al. 2018) and robustness to reward noise.

22.2 Code Repository and Open-Source Commitment

All code for the worked examples and the illustrative experiment is described in Appendix B, provided as Online Resource 2, and will be released at **[FILL: repository URL]** under **[FILL: licence]**.

22.3 Replication Checklist

The replication checklist in Table 7 applies to every benchmark and, we suggest, to every new quantum-cognition study of an autonomous system.

23 Related Surveys and Unique Contribution

Existing reviews cover neighbouring areas but not their intersection (Table 8). None of them is systematic in the PRISMA sense.

23.1 Quantum Cognition Surveys (Psychology)

The Annual Review by Pothos and Busemeyer (2022) and the overview by Huang et al. (2025b) are the key psychology surveys; Bruza et al. (2015) and Ashtiani and Azgomi (2015) are earlier reviews. They summarise the evidence from human judgement and decision-making but do not discuss robots, reinforcement learning or other autonomous agents.

Table 7 Replication checklist for quantum-cognition studies in autonomous systems

#	Requirement
1	State whether the model is quantum-like, quantum-inspired or runs on quantum hardware.
2	Pre-register the a priori QP prediction and how it could fail.
3	Report the Hilbert-space dimension, state preparation, projectors or instruments, and all free parameters.
4	Include a parameter-matched classical baseline and a noise-augmented classical baseline.
5	Evaluate on held-out data; report likelihoods and parameter counts.
6	Counterbalance question orders; test for direct influences before claiming contextuality.
7	For agents, report ≥ 30 seeds, confidence intervals and hyperparameter search budgets for all methods.
8	For hardware runs, report device, shots, noise and mitigation, and a noiseless simulation for comparison.
9	Release code, data and the order and framing of all queries.

23.2 Quantum Robotics Surveys (Including Survey 1)

Surveys of quantum robotics (Yan et al. 2024; Petschnigg et al. 2019) and our companion survey (Guru and Chauhan 2026) focus on quantum computing, sensing and communication and treat cognition briefly or not at all (Section 2.5.4).

23.3 Quantum ML and Quantum AI Surveys

Reviews of quantum machine learning (Biamonte et al. 2017; Dunjko and Briegel 2018), quantum reinforcement learning (Meyer et al. 2022; Chen 2024), quantum mathematics in AI (Widdows et al. 2021), quantum-inspired information retrieval (Uprety et al. 2020) and quantum-cognitively inspired sentiment analysis (Liu et al. 2023) cover algorithms but not the human-facing decision problems of autonomous agents. A recent chapter surveys quantum models of cognition across AI, finance and defence without grading evidence (Maksymov 2025), and position papers argue for a convergence of quantum-like cognition and AI (Khrennikov 2022, 2026).

23.4 How This Survey Differs and Adds Value

This review bridges the psychological evidence and the autonomous-systems applications and grades both; it separates quantum-like, quantum-inspired and quantum-circuit work; it appraises every application study against explicit quality criteria; and it adds reproducible case studies, an experiment, a validation roadmap and benchmarks (Section 13.4).

24 Open Problems and Future Research Directions

24.1 Near Term (1–3 Years): Validation on Real Systems

24.1.1 Comparative Studies Against Classical Baselines (Robotics, ML, DL, RL)

The first priority is controlled comparisons: QP versus parameter-matched and noise-augmented classical models, on held-out human data, in HRI and decision-support

Table 8 Coverage of this review compared with earlier reviews. ✓ = covered systematically; (✓) = covered partially; – = not covered. QC = quantum cognition; RL = reinforcement learning (quantum or quantum-inspired). The assessment is ours

Review	Focus	QC evidence and critiques	Robots and embodied agents	ML, DL and LLMs	RL	Human-AI trust	Evidence grading
Ashtiani and Azgomi (2015)	Quantum-like decision formalisms	✓	–	–	–	–	–
Bruza et al. (2015)	Quantum cognition in cognitive science	✓	–	–	–	–	–
Pothos and Busemeyer (2022)	Quantum cognition (Annual Review)	✓	–	–	–	–	–
Huang et al. (2025b)	Overview of the research programme	✓	–	–	–	(✓)	–
Dunjko and Briegel (2018)	ML and AI in the quantum domain	–	(✓)	✓	✓	–	–
Meyer et al. (2022)	Quantum reinforcement learning	–	(✓)	–	✓	–	–
Widdows et al. (2021)	Quantum mathematics in AI	(✓)	–	✓	–	–	–
Yan et al. (2024)	Quantum robotics	–	✓	(✓)	(✓)	–	–
Maksymov (2025)	Quantum models of cognition in AI applications	(✓)	(✓)	(✓)	–	(✓)	–
Guru and Chauhan (2026)	Quantum technologies for robots (companion)	–	✓	(✓)	✓	–	✓
This review	Quantum cognition for autonomous agents	✓	✓	✓	(✓)	✓	✓

settings (Benchmarks 1–4). Specific open problems are: learning QLBN phases from data with a priori constraints (Wichert et al. 2020); quantum-like models of human reward-givers for RL from human feedback (Section 10.1); quantum trust models that adapt in real time from physiological signals (Roeder et al. 2023); and QP audits of LLM components used in robots (Lo et al. 2025).

24.2 Medium Term (3–7 Years): Scaling and Optimisation

24.2.1 Efficient Classical Simulation Beyond 20 Qubits

Scaling QP models to many interacting variables requires low-rank, tensor-network or sampling-based representations ([Stoudenmire and Schwab 2016](#)), and learning algorithms that preserve the a priori constraints (such as the QQ equality) that make QP models testable.

24.2.2 Multi-Agent Coordination Using Quantum Principles

Quantum-like models of human–machine team interdependence ([Lawless et al. 2023](#)) and networks of quantum decision-makers ([Yukalov et al. 2022](#)) should be tested against classical team and opinion-dynamics models in real human–robot teams.

24.3 Long Term (7+ Years): Quantum Hardware Deployment

24.3.1 Quantum Processors for Autonomous Agent Cognition

Quantum hardware may matter if quantum agents with proven speed-ups ([Paparo et al. 2014](#); [Saggio et al. 2021](#)) can be embedded in robots, or if QP models grow beyond classical simulation. Neither is imminent.

24.3.2 Quantum-Classical Hybrid Architectures at Scale

Hybrid architectures (Figure 6) with a classical core and a quantum-like module can be deployed now; a quantum back-end is an optional long-term extension ([Widdows et al. 2023](#)).

25 Conclusion and Vision Statement

25.1 Recap: Why Quantum Cognition Matters Today Across Robotics, AI, ML, DL, RL and Physical AI

Quantum cognition offers autonomous systems a mathematically precise way to represent judgements that are sequential, ambiguous and context-dependent. The psychological evidence for such effects is substantial for order effects and interference, weaker and contested for the conjunction fallacy and contextuality, and increasingly nuanced as open-system and instrument-based models mature. This systematic review of 67 application studies finds the transfer to robots and AI promising but early: 24 are proposals or formal models, 5 reach physical systems, fewer than half compare against a classical model on the same task, and none demonstrates improved task performance in a controlled comparison.

25.2 Quantum Cognition as the Bridge to Quantum-Informed Autonomous Systems

The most defensible near-term role for quantum cognition is human-facing: predicting, explaining and aligning with human judgements, trust and preferences, where QP

models have outperformed Markov and Bayesian alternatives on several human-data tests. Physical estimation and control should stay classical. Because quantum-like models run on classical hardware, they are also the part of “quantum” autonomy that can be deployed and tested today, while the quantum-computing side reviewed in the companion survey waits for hardware (Guru and Chauhan 2026).

25.3 Call to Action: Next Steps for the Community

We call for (i) controlled comparisons against strong classical baselines; (ii) a priori, falsifiable QP predictions in every study; (iii) open benchmarks with human data collected under counterbalanced orders; (iv) clear labelling of quantum-like, quantum-inspired and quantum-hardware work; (v) released code and data; and (vi) HRI experiments that measure joint human–robot outcomes.

25.4 Vision: A Quantum-Informed AI and Robotics Future

We envisage autonomous agents that keep a classical core for perception and control but reason about people with models that respect how people actually judge. Such agents would ask questions in the right order, anticipate how their explanations change human trust, and escalate to human oversight when a person’s constructed preferences diverge from the agent’s plan. Reaching this vision depends less on quantum hardware than on careful experiments.

Declarations

- **Funding:** This research received no external funding; it was carried out as university work.
- **Competing interests:** The authors have no competing interests to declare that are relevant to the content of this article.
- **Ethics approval and consent to participate:** Not applicable. This article is a literature review; the illustrative experiment uses simulated agents only and involves no human or animal participants.
- **Consent for publication:** Not applicable.
- **Data availability:** The data-extraction table of the 67 application studies (Online Resource 3), the per-study quality appraisal (Online Resource 4), the classification of the 71 theory and evidence studies, the study summaries, the codebook and the search record are provided as supplementary material and are archived at **[FILL: Zenodo DOI]**.
- **Materials availability:** Not applicable.
- **Code availability:** The Python code that reproduces the worked examples and the illustrative reinforcement-learning experiment is provided as Online Resource 2 and described in Appendix B **[FILL: and at a public repository URL]**.
- **Author contribution:** Y.G.: conceptualisation, methodology, investigation, data curation, formal analysis, visualisation, writing – original draft. D.K.S.C.: supervision, validation, writing – review and editing. **[FILL: authors to confirm the CRediT roles]**

Supplementary information. Online Resource 1 (PDF): data-extraction table of the application studies (Table S1), PRISMA 2020 checklist (Table S2), flow of records (Table S3 and Figure S1), appraisal criteria and judgements (Tables S4 and S5), search strategy and record of the targeted search (Table S6), studies excluded under criterion E5 (Table S7), the classification of the theory and evidence studies (Table S8) and a summary of every graded study (Section S9). Online Resource 2 (Python): code for the worked examples and the illustrative experiment. Online Resource 3 (CSV): data-extraction table. Online Resource 4 (CSV): quality appraisal.

Appendix A Mathematical Formalism (Proofs and Derivations)

A.1 Born and Lüders rules

For a pure state $|\psi\rangle$ and projector P , $p = \|P|\psi\rangle\|^2 = \langle\psi|P|\psi\rangle$, and the post-measurement state is $P|\psi\rangle/\|P|\psi\rangle\|$. For a density matrix ρ , $p = \text{tr}(P\rho)$ and $\rho \mapsto P\rho P/\text{tr}(P\rho)$.

A.2 Sequential probabilities and the order effect

Because P is idempotent and self-adjoint, $p(A \text{ then } B) = \|P_B P_A \psi\|^2 = \langle\psi|P_A P_B P_A|\psi\rangle$. Hence the order effect is $p(A \text{ then } B) - p(B \text{ then } A) = \langle\psi|P_A P_B P_A - P_B P_A P_B|\psi\rangle$, which vanishes for all ψ when $[P_A, P_B] = 0$.

A.3 Proof of the QQ equality

Write $\langle X \rangle = \langle\psi|X|\psi\rangle$ and $\bar{A} = I - P_A$. Then $p(A_y B_n) = \langle P_A \rangle - \langle P_A P_B P_A \rangle$ and $p(A_n B_y) = \langle \bar{A} P_B \bar{A} \rangle = \langle P_B \rangle - \langle P_A P_B \rangle - \langle P_B P_A \rangle + \langle P_A P_B P_A \rangle$. Their sum is $\langle P_A \rangle + \langle P_B \rangle - \langle P_A P_B \rangle - \langle P_B P_A \rangle$, which is symmetric in A and B . The same sum for the reverse order is therefore identical, and $[p(A_y B_n) + p(A_n B_y)] - [p(B_y A_n) + p(B_n A_y)] = 0$ for any pure state and any pair of projectors (Wang and Busemeyer 2013).

A.4 Interference and the law of total probability

Let the unresolved state be $|\psi\rangle = a|\psi_1\rangle + b|\psi_2\rangle$, where $|\psi_1\rangle$ and $|\psi_2\rangle$ are the states after learning each resolution, and $|a|^2 + |b|^2 = 1$. Then $p(D) = \|P_D \psi\|^2 = |a|^2 p_1 + |b|^2 p_2 + 2 \text{Re}(a^* b \langle\psi_1|P_D|\psi_2\rangle)$, where $p_i = \|P_D \psi_i\|^2$. Writing the cross term as $2\sqrt{|a|^2 p_1 |b|^2 p_2} \cos \theta$ gives the interference form used in Section 5.5. With $p_1 = 0.97$, $p_2 = 0.84$ and $|a|^2 = |b|^2 = 0.5$, the observed 0.63 implies $\cos \theta = -0.305$ ($\theta \approx 107.7^\circ$).

A.5 Condition for a conjunction fallacy

In a real two-dimensional space with the state at angle 0, “feminist” at α and “bank teller” at β , $p(\text{BT}) = \cos^2 \beta$ and $p(\text{F then BT}) = \cos^2 \alpha \cos^2(\beta - \alpha)$. The fallacy $p(\text{F then BT}) > p(\text{BT})$ occurs whenever $\cos \alpha \cos(\beta - \alpha) > \cos \beta$, which holds, for example, for $\alpha = 40^\circ$ and $\beta = 80^\circ$ ($0.344 > 0.030$).

A.6 Open-system (Lindblad) dynamics

Quantum–Markov models evolve a density matrix according to $d\rho/dt = -i[H, \rho] + \sum_k \gamma_k \left(L_k \rho L_k^\dagger - \frac{1}{2} \{L_k^\dagger L_k, \rho\} \right)$, where the Hamiltonian H generates coherent (quantum-walk) deliberation and the operators L_k generate incoherent (Markov-walk) transitions (Busemeyer et al. 2020; Asano and Khrennikov 2026). Setting $\gamma_k = 0$ recovers unitary dynamics; setting $H = 0$ with suitable L_k recovers a Markov process.

A.7 Quantum instruments

A quantum instrument assigns to each outcome a completely positive map that may differ from the Lüders projection. Instruments allow response replicability and order effects to coexist (Ozawa and Khrennikov 2020, 2021) and underlie recent characterisations of cognitive measurements (Fuyama et al. 2025).

Appendix B Implementation Code and Framework (Python, NumPy; Qiskit, PennyLane, TensorFlow/PyTorch Integration)

Online Resource 2 is a single NumPy file that reproduces every computed value of the review: the order effect and QQ equality (Worked Example 1), the conjunction fallacy (Example 2), the disjunction-effect interference term (Example 3), Kalman and two-cue QLBN fusion (Case Study 1), order-dependent labelling (Case Study 2) and the reinforcement-learning experiment (Case Study 3, run with the `--r1` option). Its core is two functions:

```
import numpy as np

def proj(theta):
    """Rank-1 projector onto the unit vector at angle theta in a real 2-D Hilbert space."""
    v = np.array([np.cos(theta), np.sin(theta)])
    return np.outer(v, v)

def seq_prob(psi, P1, P2):
    """Lueders rule: probability of 'yes' to question 1 then 'yes' to question 2."""
    return float(np.linalg.norm(P2 @ P1 @ psi) ** 2)

# Worked Example 1: state at 20 deg, A at 0 deg, B at 55 deg
psi = np.array([np.cos(np.deg2rad(20)), np.sin(np.deg2rad(20))])
A, B = proj(0.0), proj(np.deg2rad(55))
print(seq_prob(psi, A, B), seq_prob(psi, B, A)) # 0.291 0.221
```

The same projectors can be written as circuits for Qiskit or PennyLane (Bergholm et al. 2018), with a state-preparation rotation and measurements in rotated bases, and a QP module of this size can be wrapped as a non-trainable or trainable layer in PyTorch or TensorFlow (Broughton et al. 2020); for the dimensions used in cognitive models a NumPy implementation is faster than either (Section 20.3).

Appendix C Experimental Data and Raw Results

Table C1 reports the illustrative reinforcement-learning experiment of Case Study 3 (Section 16.3) and Table C2 the values of the worked examples and case studies. No other experiments were run for this review. All values can be reproduced with the code in Online Resource 2.

Table C1 Case Study 3 (Section 16.3): 6×6 grid, 300 episodes, 30 seeds. Final = mean steps over the last 50 episodes (optimum 10); convergence = first episode at which the 10-episode moving average is ≤ 15 steps

Agent	Reward noise σ	Final steps (mean $\pm$ SD)	Convergence episode (mean $\pm$ SD)	Seeds converged	Total steps (mean)
ϵ -greedy Q-learning	0	10.1 ± 0.0	60 ± 6	30/30	4,453
Softmax Q-learning	0	10.0 ± 0.0	128 ± 10	30/30	8,163
Quantum-inspired RL ($k = 30$)	0	10.0 ± 0.0	134 ± 20	30/30	7,362
ϵ -greedy Q-learning	0.5	13.8 ± 13.4	154 ± 41	30/30	24,746
Softmax Q-learning	0.5	16.2 ± 16.7	187 ± 28	30/30	32,885
Quantum-inspired RL ($k = 30$)	0.5	67.1 ± 40.1	132 ± 42	25/30	25,122

Sensitivity of quantum-inspired RL to the amplification constant k (30 seeds; mean steps over the last 50 episodes / seeds converged): $k = 10$: $16.3 / 30$ ($\sigma = 0$) and $78.1 / 5$ ($\sigma = 0.5$); $k = 30$: $10.0 / 30$ and $67.1 / 25$; $k = 100$: $10.0 / 30$ and $73.5 / 25$.

Table C2 Values of the worked examples and case studies (computed with the code in Online Resource 2)

Example	Setup	Result
1 Order effect (Section 5.3)	State 20° , A 0° , B 55°	$p(A) = 0.883$; $p(B) = 0.671$; $p(A \rightarrow B) = 0.291$; $p(B \rightarrow A) = 0.221$; QQ = 0 $p(BT) = 0.030$; $p(F \rightarrow BT) = 0.344$
2 Conjunction fallacy (Section 5.4)	State 0° , feminist 40° , bank teller 80°	
3 Interference (Section 5.5)	Shafir and Tversky (1992): 0.97, 0.84, 0.63	Classical 0.905; interference -0.275 ; $\theta = 107.7^\circ$
Case Study 1: Kalman fusion (Section 16.1)	$z_1 = 2.10$ (var 0.04), $z_2 = 1.90$ (var 0.09)	$\hat{x} = 2.038$, var 0.028; order-invariant
Case Study 1: two-cue QLBN	$p_1 = 0.80$, $p_2 = 0.30$, equal weights	0.795 (60°), 0.55 (90°), 0.305 (120°)
Case Study 2: label order (Section 16.2)	Image 35° , cat 0° , wild 60°	$p(\text{cat} \rightarrow \text{wild}) = 0.168$; $p(\text{wild} \rightarrow \text{cat}) = 0.205$

Appendix D Benchmark Protocols and Reproducibility Checklist

Each benchmark of Section 22 should be run with a pre-registered protocol: (1) state the QP model and its a priori prediction (for example the QQ equality or a total-probability violation) before collecting data; (2) collect human judgements under counterbalanced question orders, with a sample size set by a power analysis; (3) fit the QP model, a parameter-matched classical model and a noise-augmented classical model on training data; (4) compare predictive performance on held-out data with cross-validated log-likelihood and report parameter counts; (5) for agent benchmarks, report task outcomes over 30 or more seeds with confidence intervals (Henderson et al. 2018); and (6) release code, data and the order and framing of all queries. The replication checklist of Table 7 lists the items a report must contain.

Appendix E Extended Case Studies and Worked Examples

A question-order test for human–robot interaction

Suppose a service robot asks a user two questions before acting: “Is the task urgent?” (A) and “Should I follow the usual route?” (B). With the belief geometry of Worked Example 1, the model predicts $p(A \text{ yes, then } B \text{ yes}) = 0.29$ and $p(B \text{ yes, then } A \text{ yes}) = 0.22$, and the QQ equality constrains the four sequential patterns. A counterbalanced study, with a sample size set by a power analysis, can test the prediction; if it holds, the robot can choose the order that minimises misunderstanding, and log the order for accountability (Section 19.3).

A two-cue judgement with a quantum-like Bayesian network

Case Study 1 computed the interference term for two cues. In a full QLBN with n unobserved binary nodes, there are 2^n paths and up to $2^n(2^n - 1)/2$ pairwise interference terms. Heuristics that set all phases from a single principle keep the model testable (Moreira and Wichert 2016; Wichert et al. 2020). The renormalisation step (dividing by the sum over outcomes) must always be applied, because unnormalised quantum-like probabilities can exceed 1, as $\theta = 0$ in Case Study 1 shows.

Adding de-amplification to quantum-inspired RL

The instability found in Case Study 3 suggests a simple modification: when the TD error is negative, apply a rotation that moves amplitude away from the chosen action, for example a reflection about the uniform state. Testing this against ε -greedy and softmax under the Benchmark 4 protocol is a natural next experiment.

References

Aaronson S (2015) Read the fine print. Nature Physics 11(4):291–293. <https://doi.org/10.1038/nphys3272>

- Acuto A, Barillà P, Bozzolo L, et al (2022) Variational quantum soft actor-critic for robotic arm control. [arXiv:2212.11681](#)
- Aerts D (2009) Quantum structure in cognition. *Journal of Mathematical Psychology* 53(5):314–348. <https://doi.org/10.1016/j.jmp.2009.04.005>
- Aerts D, Sassoli de Bianchi M (2017) Beyond-quantum modeling of question order effects and response replicability in psychological measurements. *Journal of Mathematical Psychology* 79:104–120. [arXiv:1508.03686](#)
- Aerts D, Sozzo S (2011) Quantum structure in cognition: Why and how concepts are entangled. In: Song D, Melucci M, Frommholz I, et al (eds) *Quantum Interaction: 5th International Symposium, QI 2011, Aberdeen, UK, June 26–29, 2011, Revised Selected Papers, Lecture Notes in Computer Science*, vol 7052. Springer, Berlin, Heidelberg, pp 116–127, https://doi.org/10.1007/978-3-642-24971-6_12
- Aerts D, D’Hondt E, Gabora L (2000) Why the disjunction in quantum logic is not classical. *Foundations of Physics* 30(9):1473–1480. <https://doi.org/10.1023/A:1026457918361>
- Aerts D, Broekaert J, Czachor M, et al (2011a) A quantum-conceptual explanation of violations of expected utility in economics. In: Song D, Melucci M, Frommholz I, et al (eds) *Quantum Interaction: 5th International Symposium, QI 2011, Aberdeen, UK, June 26–29, 2011, Revised Selected Papers, Lecture Notes in Computer Science*, vol 7052. Springer, Berlin, Heidelberg, pp 192–198, https://doi.org/10.1007/978-3-642-24971-6_19
- Aerts D, Czachor M, Sozzo S (2011b) Quantum interaction approach in cognition, artificial intelligence and robotics. In: Privman V, Ovchinnikov V (eds) *Proceedings of the Fifth International Conference on Quantum, Nano and Micro Technologies (ICQNM 2011)*. IARIA, pp 35–40, [arXiv:1104.3345](#)
- Aerts D, Gabora L, Sozzo S, et al (2011c) Quantum structure in cognition: Fundamentals and applications. In: *Proceedings of the Fifth International Conference on Quantum, Nano and Micro Technologies (ICQNM 2011)*. IARIA, [arXiv:1104.3344](#)
- Aerts D, Gabora L, Sozzo S (2013) Concepts and their dynamics: A quantum-theoretic modeling of human thought. *Topics in Cognitive Science* 5(4):737–772. <https://doi.org/10.1111/tops.12042>
- Aerts D, Aerts Arguëlles J, Beltran L, et al (2017) Testing quantum models of conjunction fallacy on the World Wide Web. *International Journal of Theoretical Physics* 56(12):3744–3756. <https://doi.org/10.1007/s10773-017-3288-8>
- Aerts D, Aerts Arguëlles J, Beltran L, et al (2026) Identifying quantum structure in AI language: Evidence for evolutionary convergence of human and artificial cognition. *Entropy* 28(6):622. <https://doi.org/10.3390/e28060622>

- Asano M, Khrennikov A (2026) Quantum-like models of cognition and decision making: open-systems and Gorini–Kossakowski–Sudarshan–Lindblad dynamics. [arXiv:2604.18643](#)
- Ashtiani M, Abdollahi Azgomi M (2019) A novel trust evolution algorithm based on a quantum-like model of computational trust. *Cognition, Technology & Work* 21(2):201–224. <https://doi.org/10.1007/s10111-018-0496-9>
- Ashtiani M, Azgomi MA (2015) A survey of quantum-like approaches to decision making and cognition. *Mathematical Social Sciences* 75:49–80. <https://doi.org/10.1016/j.mathsocsci.2015.02.004>
- Basieva I, Cervantes VH, Dzhamfarov EN, et al (2019) True contextuality beats direct influences in human decision making. *Journal of Experimental Psychology: General* 148(11):1925–1937. <https://doi.org/10.1037/xge0000585>
- Benioff P (1998) Quantum robots and environments. *Physical Review A* 58(2):893–904. <https://doi.org/10.1103/PhysRevA.58.893>
- Bergholm V, Izaac J, Schuld M, et al (2018) PennyLane: Automatic differentiation of hybrid quantum-classical computations. [arXiv:1811.04968](#)
- Berner N, Fortuin V, Landman J (2021) Quantum Bayesian neural networks. <https://doi.org/10.48550/arXiv.2107.09599>, [arXiv:2107.09599](#)
- Biamonte J, Wittek P, Pancotti N, et al (2017) Quantum machine learning. *Nature* 549(7671):195–202. <https://doi.org/10.1038/nature23474>
- Binz M, Schulz E (2023) Using cognitive psychology to understand GPT-3. *Proceedings of the National Academy of Sciences* 120(6):e2218523120. <https://doi.org/10.1073/pnas.2218523120>
- Blanke M, Kinnaert M, Lunze J, et al (2016) *Diagnosis and Fault-Tolerant Control*, 3rd edn. Springer, Berlin, Heidelberg, <https://doi.org/10.1007/978-3-662-47943-8>
- Blutner R, Pothos EM, Bruza P (2013) A quantum probability perspective on borderline vagueness. *Topics in Cognitive Science* 5(4):711–736. <https://doi.org/10.1111/tops.12041>
- Boyajian WL, Clausen J, Trenkwalder LM, et al (2020) On the convergence of projective-simulation-based reinforcement learning in Markov decision processes. *Quantum Machine Intelligence* 2(2):13. <https://doi.org/10.1007/s42484-020-00023-9>
- Boyer-Kassem T, Duchêne S, Guerci E (2016) Quantum-like models cannot account for the conjunction fallacy. *Theory and Decision* 81(4):479–510. <https://doi.org/10.1007/s11238-016-9549-9>

- Briegel HJ, De las Cuevas G (2012) Projective simulation for artificial intelligence. *Scientific Reports* 2:400. <https://doi.org/10.1038/srep00400>
- Broekaert J, Basieva I, Blasiak P, et al (2017) Quantum-like dynamics applied to cognition: A consideration of available options. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences* 375(2106):20160387. <https://doi.org/10.1098/rsta.2016.0387>
- Broughton M, Verdon G, McCourt T, et al (2020) TensorFlow Quantum: A software framework for quantum machine learning. [arXiv:2003.02989](https://arxiv.org/abs/2003.02989)
- Bruza P, Kitto K, Nelson D, et al (2009) Is there something quantum-like about the human mental lexicon? *Journal of Mathematical Psychology* 53(5):363–377. <https://doi.org/10.1016/j.jmp.2009.04.004>
- Bruza PD, Wang Z, Busemeyer JR (2015) Quantum cognition: A new theoretical approach to psychology. *Trends in Cognitive Sciences* 19(7):383–393. <https://doi.org/10.1016/j.tics.2015.05.001>
- Busemeyer JR, Bruza PD (2012) *Quantum Models of Cognition and Decision*. Cambridge University Press, Cambridge, <https://doi.org/10.1017/CBO9780511997716>
- Busemeyer JR, Wang Z (2015) What is quantum cognition, and how is it applied to psychology? *Current Directions in Psychological Science* 24(3):163–169. <https://doi.org/10.1177/0963721414568663>
- Busemeyer JR, Wang Z (2017) Is there a problem with quantum models of psychological measurements? *PLoS ONE* 12(11):e0187733. <https://doi.org/10.1371/journal.pone.0187733>
- Busemeyer JR, Wang Z (2018) Hilbert space multidimensional theory. *Psychological Review* 125(4):572–591. <https://doi.org/10.1037/rev0000106>
- Busemeyer JR, Wang Z, Townsend JT (2006) Quantum dynamics of human decision-making. *Journal of Mathematical Psychology* 50(3):220–241. <https://doi.org/10.1016/j.jmp.2006.01.003>
- Busemeyer JR, Wang Z, Lambert-Mogiliansky A (2009) Empirical comparison of Markov and quantum models of decision making. *Journal of Mathematical Psychology* 53(5):423–433. <https://doi.org/10.1016/j.jmp.2009.03.002>
- Busemeyer JR, Pothos EM, Franco R, et al (2011) A quantum theoretical explanation for probability judgment errors. *Psychological Review* 118(2):193–218. <https://doi.org/10.1037/a0022542>

- Busemeyer JR, Fakhari P, Kvam P (2017) Neural implementation of operations used in quantum cognition. *Progress in Biophysics and Molecular Biology* 130(Part A):53–60. <https://doi.org/10.1016/j.pbiomolbio.2017.04.007>
- Busemeyer JR, Kvam PD, Pleskac TJ (2019) Markov versus quantum dynamic models of belief change during evidence monitoring. *Scientific Reports* 9:18025. <https://doi.org/10.1038/s41598-019-54383-9>
- Busemeyer JR, Zhang Q, Balakrishnan SN, et al (2020) Application of quantum–Markov open system models to human cognition and decision. *Entropy* 22(9):990. <https://doi.org/10.3390/e22090990>
- Busemeyer JR, Ozawa M, Pothos EM, et al (2025) Incorporating episodic memory into quantum models of judgment and decision. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences* 383(2309):20240387. <https://doi.org/10.1098/rsta.2024.0387>
- Cerezo M, Sone A, Volkoff T, et al (2021) Cost function dependent barren plateaus in shallow parametrized quantum circuits. *Nature Communications* 12:1791. <https://doi.org/10.1038/s41467-021-21728-w>
- Cervantes VH, Dzhamfarov EN (2018) Snow Queen is evil and beautiful: Experimental evidence for probabilistic contextuality in human choices. *Decision* 5(3):193–204. <https://doi.org/10.1037/dec0000095>
- Chella A, Gaglio S, Pilato G, et al (2022) A quantum planner for robot motion. *Mathematics* 10(14):2475. <https://doi.org/10.3390/math10142475>
- Chella A, Gaglio S, Pilato G, et al (2026) An architecture for a quantum teleo-reactive robot. *Entropy* 28(7):731. <https://doi.org/10.3390/e28070731>
- Chen SYC (2024) An introduction to quantum reinforcement learning (QRL). [arXiv:2409.05846](https://arxiv.org/abs/2409.05846)
- Chen SYC, Yang CHH, Qi J, et al (2020) Variational quantum circuits for deep reinforcement learning. *IEEE Access* 8:141007–141024. <https://doi.org/10.1109/ACCESS.2020.3010470>
- Clark A (2013) Whatever next? predictive brains, situated agents, and the future of cognitive science. *Behavioral and Brain Sciences* 36(3):181–204. <https://doi.org/10.1017/S0140525X12000477>
- Costello F, Watts P (2014) Surprisingly rational: probability theory plus noise explains biases in judgment. *Psychological Review* 121(3):463–480. <https://doi.org/10.1037/a0037010>

- Costello F, Watts P (2018) Invariants in probabilistic reasoning. *Cognitive Psychology* 100:1–16. <https://doi.org/10.1016/j.cogpsych.2017.11.003>
- Daglarli E (2025) A generative neuro-cognitive architecture using quantum algorithms for the autonomous behavior of a smart agent in a simulation environment. *Computers, Materials & Continua* 84(3). <https://doi.org/10.32604/cmc.2025.065572>
- De Carolis BN, Loglisci C, Palestra G, et al (2025) Quantum-enhanced social robotics: The QUADRI project. In: *Adjunct Proceedings of the 33rd ACM Conference on User Modeling, Adaptation and Personalization (UMAP Adjunct '25)*. ACM, New York, NY, USA, <https://doi.org/10.1145/3708319.3735545>
- Dong D, Chen C, Zhang C, et al (2006) Quantum robot: Structure, algorithms and applications. *Robotica* 24(4):513–521. <https://doi.org/10.1017/S0263574705002596>
- Dong D, Chen C, Li H, et al (2008) Quantum reinforcement learning. *IEEE Transactions on Systems, Man, and Cybernetics, Part B: Cybernetics* 38(5):1207–1220. <https://doi.org/10.1109/TSMCB.2008.925743>
- Dong D, Chen C, Chu J, et al (2012) Robust quantum-inspired reinforcement learning for robot navigation. *IEEE/ASME Transactions on Mechatronics* 17(1):86–97. <https://doi.org/10.1109/TMECH.2010.2090896>
- Drossi T, Fremont DJ, Ghosh S, et al (2019) VerifAI: A toolkit for the formal design and analysis of artificial intelligence-based systems. In: Dillig I, Tasiran S (eds) *Computer Aided Verification (CAV 2019)*, *Lecture Notes in Computer Science*, vol 11561. Springer, Cham, pp 432–442, https://doi.org/10.1007/978-3-030-25540-4_25
- Dunjko V, Briegel HJ (2018) Machine learning & artificial intelligence in the quantum domain: a review of recent progress. *Reports on Progress in Physics* 81(7):074001. <https://doi.org/10.1088/1361-6633/aab406>
- Dzhafarov EN, Kujala JV (2016) Context–content systems of random variables: The Contextuality-by-Default theory. *Journal of Mathematical Psychology* 74:11–33. <https://doi.org/10.1016/j.jmp.2016.04.010>
- Dzhafarov EN, Zhang R, Kujala J (2016) Is there contextuality in behavioural and social systems? *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences* 374(2058):20150099. <https://doi.org/10.1098/rsta.2015.0099>
- Eisert J, Wilkens M, Lewenstein M (1999) Quantum games and quantum strategies. *Physical Review Letters* 83(15):3077–3080. <https://doi.org/10.1103/PhysRevLett.83.3077>
- Ellsberg D (1961) Risk, ambiguity, and the Savage axioms. *The Quarterly Journal of Economics* 75(4):643–669. <https://doi.org/10.2307/1884324>

- Epping GP, Kvam PD, Pleskac TJ, et al (2023) Open system model of choice and response time. *Journal of Choice Modelling* 49:100453. <https://doi.org/10.1016/j.jocm.2023.100453>
- Essalmi K, Garrido F, Nashashibi F (2025) Quantum game models for interaction-aware decision-making in automated driving. ArXiv preprint, [arXiv:2509.01582](https://arxiv.org/abs/2509.01582)
- Essalmi K, Garrido F, Nashashibi F (2026) Multi-player, multi-strategy quantum game model for interaction-aware decision-making in automated driving. ArXiv preprint, [arXiv:2602.03571](https://arxiv.org/abs/2602.03571)
- Fisher MPA (2015) Quantum cognition: The possibility of processing with nuclear spins in the brain. *Annals of Physics* 362:593–602. <https://doi.org/10.1016/j.aop.2015.08.020>
- Flamini F, Hamann A, Jerbi S, et al (2020) Photonic architecture for reinforcement learning. *New Journal of Physics* 22(4):045002. <https://doi.org/10.1088/1367-2630/ab783c>
- Flamini F, Krumm M, Fiderer LJ, et al (2024) Towards interpretable quantum machine learning via single-photon quantum walks. *Quantum Science and Technology* 9(4):045011. <https://doi.org/10.1088/2058-9565/ad5907>
- Friston K (2010) The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience* 11(2):127–138. <https://doi.org/10.1038/nrn2787>
- Fuyama M, Khrennikov A, Ozawa M (2025) Quantum-like cognition and decision-making in the light of quantum measurement theory. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences* 383(2309):20240372. <https://doi.org/10.1098/rsta.2024.0372>
- Gao Q, Dong D, Petersen IR, et al (2021) Design of a discrete-time fault-tolerant quantum filter and fault detector. *IEEE Transactions on Cybernetics* 51(2):889–899. <https://doi.org/10.1109/TCYB.2019.2899877>
- Gao X, Yang H, Pan L, et al (2025) Quantum-like evidence networks decision-making model. *Engineering Applications of Artificial Intelligence* 157:111368. <https://doi.org/10.1016/j.engappai.2025.111368>
- Gkoumas D, Li Q, Dehdashti S, et al (2021) Quantum cognitively motivated decision fusion for video sentiment analysis. In: *Proceedings of the AAAI Conference on Artificial Intelligence*, pp 827–835, <https://doi.org/10.1609/aaai.v35i1.16165>
- Guru Y, Chauhan DKS (2026) Quantum technologies for autonomous robotics: A systematic survey of methods, evidence, and deployment constraints. Companion manuscript, prepared for *ACM Computing Surveys*

- Hameroff S, Penrose R (2014) Consciousness in the universe: A review of the ‘Orch OR’ theory. *Physics of Life Reviews* 11(1):39–78. <https://doi.org/10.1016/j.plrev.2013.08.002>
- Hangl S, Ugur E, Szedmak S, et al (2016) Robotic playing for hierarchical complex skill learning. In: *Proceedings of the 2016 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. IEEE, pp 2799–2804, [arXiv:1603.00794](https://arxiv.org/abs/1603.00794)
- Hangl S, Dunjko V, Briegel HJ, et al (2020) Skill learning by autonomous robotic playing using active learning and exploratory behavior composition. *Frontiers in Robotics and AI* 7:42. <https://doi.org/10.3389/frobt.2020.00042>
- Haykin S (ed) (2001) *Kalman Filtering and Neural Networks*. Wiley, New York, <https://doi.org/10.1002/0471221546>
- Henderson P, Islam R, Bachman P, et al (2018) Deep reinforcement learning that matters. In: *Proceedings of the AAAI Conference on Artificial Intelligence*, pp 3207–3214, <https://doi.org/10.1609/aaai.v32i1.11694>
- Ho JKW, Hoorn JF (2022) Quantum affective processes for multidimensional decision-making. *Scientific Reports* 12:20468. <https://doi.org/10.1038/s41598-022-22855-0>
- Hohenfeld H, Heimann D, Wiebe F, et al (2024) Quantum deep reinforcement learning for robot navigation tasks. *IEEE Access* 12:87217–87236. <https://doi.org/10.1109/ACCESS.2024.3417808>
- Hoorn JF, Ho JKW (2019) Robot affect: the amygdala as Bloch sphere. *ArXiv preprint*, [arXiv:1911.12128](https://arxiv.org/abs/1911.12128)
- Huang J, Busemeyer JR, Ebel Z, et al (2025a) Bridging the gap between subjective probability and probability judgments: the Quantum Sequential Sampler. *Psychological Review* 132(4):916–955. <https://doi.org/10.1037/rev0000489>
- Huang J, Epping GP, Trueblood JS, et al (2025b) An overview of the quantum cognition research program. *Psychonomic Bulletin & Review* 32:2507–2556. <https://doi.org/10.3758/s13423-025-02675-9>
- Huang J, Pleskac TJ, Busemeyer JR (2025c) Quantum recognition heuristics: A tutorial example of quantum heuristics. In: *2025 IEEE International Conference on Quantum Computing and Engineering (QCE)*. IEEE, <https://doi.org/10.1109/QCE65121.2025.10486>
- Humr S, Canan M (2026) System level intervention for AI-supported decision-making: A quantum cognition perspective. In: *Quantum AI and NLP (QNLPAI 2025)*, *Communications in Computer and Information Science*, vol 2659. Springer, Cham, p 250–273, https://doi.org/10.1007/978-3-032-13883-5_14

- Humr S, Canan M, Demir M (2025) A quantum probability approach to improving human–AI decision making. *Entropy* 27(2):152. <https://doi.org/10.3390/e27020152>
- Humr SA, Canan M, Demir M (2023) Temporal evolution of trust in artificial intelligence-supported decision-making. *Proceedings of the Human Factors and Ergonomics Society Annual Meeting* 67(1):1936–1941. <https://doi.org/10.1177/21695067231193672>
- Imannezhad P, Pothos EM, Wills AJ (2026) Divergent patterns of probabilistic reasoning in humans and GPT-5. *Frontiers in Psychology* 17:1782184. <https://doi.org/10.3389/fpsyg.2026.1782184>
- Jerbi S, Gyurik C, Marshall SC, et al (2021) Parametrized quantum policies for reinforcement learning. In: *Advances in Neural Information Processing Systems*, [arXiv:2103.05577](https://arxiv.org/abs/2103.05577)
- Kak S (2018) Order effects for queries in intelligent systems. *ArXiv preprint*, [arXiv:1804.02759](https://arxiv.org/abs/1804.02759)
- Kalman RE (1960) A new approach to linear filtering and prediction problems. *Journal of Basic Engineering* 82(1):35–45. <https://doi.org/10.1115/1.3662552>
- Kang P (2026) Auditing question-order effects in large language models with the QQ equality: Mechanism characterization and a saturation caveat. [arXiv:2607.17219](https://arxiv.org/abs/2607.17219)
- Kellen D, Singmann H, Batchelder WH (2018) Classic-probability accounts of mirrored (quantum-like) order effects in human judgments. *Decision* 5(4):323–338. <https://doi.org/10.1037/dec0000080>
- Khoshnoud F, Quadrelli MB, Esat II, et al (2020) Quantum cooperative robotics and autonomy. [arXiv:2008.12230](https://arxiv.org/abs/2008.12230)
- Khrennikov A (1999) Classical and quantum mechanics on information spaces with applications to cognitive, psychological, social, and anomalous phenomena. *Foundations of Physics* 29(7):1065–1098. <https://doi.org/10.1023/A:1018885632116>
- Khrennikov A (2010) *Ubiquitous Quantum Structure: From Psychology to Finance*. Springer, Berlin, Heidelberg, <https://doi.org/10.1007/978-3-642-05101-2>
- Khrennikov A (2022) Quantum probability for modeling cognition, decision making, and artificial intelligence. In: Accardi L, Mukhamedov F, Al Rawashdeh A (eds) *Infinite Dimensional Analysis, Quantum Probability and Applications, Springer Proceedings in Mathematics & Statistics*, vol 390. Springer, Cham, p 89–100, https://doi.org/10.1007/978-3-031-06170-7_4
- Khrennikov A (2023) Open systems, quantum probability, and logic for quantum-like modeling in biology, cognition, and decision-making. *Entropy* 25(6):886. <https://doi.org/10.3390/entropy25060886>

[//doi.org/10.3390/e25060886](https://doi.org/10.3390/e25060886)

- Khrennikov A (2024) Contextual measurement model and quantum theory. *Royal Society Open Science* 11(3):231953. <https://doi.org/10.1098/rsos.231953>
- Khrennikov A (2026) Quantum-like cognition and AI toward the convergence of natural and artificial intelligence. *Discover Artificial Intelligence* 6:167. <https://doi.org/10.1007/s44163-026-00909-w>
- Khrennikov A, Yamada M (2025) Quantum-like representation of neuronal networks' activity: modeling "mental entanglement". *Frontiers in Human Neuroscience* 19:1685339. <https://doi.org/10.3389/fnhum.2025.1685339>
- Khrennikov A, Basieva I, Dzhaferov EN, et al (2014) Quantum models for psychological measurements: An unsolved problem. *PLoS ONE* 9(10):e110909. <https://doi.org/10.1371/journal.pone.0110909>
- Khrennikov A, Basieva I, Pothos EM, et al (2018) Quantum probability in decision making from quantum information representation of neuronal states. *Scientific Reports* 8:16225. <https://doi.org/10.1038/s41598-018-34531-3>
- Khrennikov A, Iriki A, Basieva I (2025a) Constructing a bridge between functioning of oscillatory neuronal networks and quantum-like cognition along with quantum-inspired computation and AI. *BioSystems* 257:105573. <https://doi.org/10.1016/j.biosystems.2025.105573>
- Khrennikov A, Ozawa M, Benninger F, et al (2025b) Coupling quantum-like cognition with the neuronal networks within generalized probability theory. *Journal of Mathematical Psychology* 125:102923. <https://doi.org/10.1016/j.jmp.2025.102923>
- Kitchenham B, Charters S (2007) Guidelines for performing systematic literature reviews in software engineering. EBSE Technical Report EBSE-2007-01, Keele University and Durham University, version 2.3
- Kong D, Xue K, Wang P, et al (2026) Emergence of longitudinal queues in group navigation: An interpretable approach via projective simulation. *Biomimetics* 11(3):201. <https://doi.org/10.3390/biomimetics11030201>
- Kotseruba I, Tsotsos JK (2020) 40 years of cognitive architectures: Core cognitive abilities and practical applications. *Artificial Intelligence Review* 53(1):17–94. <https://doi.org/10.1007/s10462-018-9646-y>
- Krizhevsky A, Sutskever I, Hinton GE (2012) ImageNet classification with deep convolutional neural networks. In: Pereira F, Burges CJC, Bottou L, et al (eds) *Advances in Neural Information Processing Systems* 25 (NIPS 2012). Curran Associates, pp 1097–1105, URL <https://papers.nips.cc/paper/2012/hash/c399862d3b9d6b76c8436e924a68c45b-Abstract.html>

- Kvam PD, Pleskac TJ, Yu S, et al (2015) Interference effects of choice on confidence: Quantum characteristics of evidence accumulation. *Proceedings of the National Academy of Sciences* 112(34):10645–10650. <https://doi.org/10.1073/pnas.1500688112>
- Kvam PD, Busemeyer JR, Pleskac TJ (2021) Temporal oscillations in preference strength provide evidence for an open system model of constructed preference. *Scientific Reports* 11:8169. <https://doi.org/10.1038/s41598-021-87659-0>
- Laird JE, Lebiere C, Rosenbloom PS (2017) A standard model of the mind: Toward a common computational framework across artificial intelligence, cognitive science, neuroscience, and robotics. *AI Magazine* 38(4):13–26. <https://doi.org/10.1609/aimag.v38i4.2744>
- Lanza D, Solinas P, Mastrogiovanni F (2020) A preliminary study for a quantum-like robot perception model. [arXiv:2006.02771](https://arxiv.org/abs/2006.02771)
- Lanza D, Solinas P, Mastrogiovanni F (2021) Multi-sensory integration in a quantum-like robot perception model. In: Siciliano B, Laschi C, Khatib O (eds) *Experimental Robotics: The 17th International Symposium (ISER 2020)*, Springer Proceedings in Advanced Robotics, vol 19. Springer, Cham, p 502–509, https://doi.org/10.1007/978-3-030-71151-1_44
- Lauri M, Hsu D, Pajarinen J (2023) Partially observable Markov decision processes in robotics: a survey. *IEEE Transactions on Robotics* 39(1):21–40. <https://doi.org/10.1109/TRO.2022.3200138>
- Lawless WF (2020) Quantum-like interdependence theory advances autonomous human–machine teams (A-HMTs). *Entropy* 22(11):1227. <https://doi.org/10.3390/e22111227>
- Lawless WF, Moskowitz IS (2025) An entropy approach to interdependent human–machine teams. *Entropy* 27(2):176. <https://doi.org/10.3390/e27020176>
- Lawless WF, Moskowitz IS, Doctor KZ (2023) A quantum-like model of interdependence for embodied human–machine teams: Reviewing the path to autonomy facing complexity and uncertainty. *Entropy* 25(9):1323. <https://doi.org/10.3390/e25091323>
- Li Q, Wang B, Melucci M (2019) CNM: An interpretable complex-valued network for matching. In: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*. Association for Computational Linguistics, Minneapolis, Minnesota, pp 4139–4148, <https://doi.org/10.18653/v1/N19-1420>
- Li Q, Gkoumas D, Lioma C, et al (2021a) Quantum-inspired multimodal fusion for video sentiment analysis. *Information Fusion* 65:58–71. <https://doi.org/10.1016/j.inffus.2021.05.001>

- Li Q, Gkoumas D, Sordoni A, et al (2021b) Quantum-inspired neural network for conversational emotion recognition. In: Proceedings of the AAAI Conference on Artificial Intelligence, pp 13270–13278. <https://doi.org/10.1609/aaai.v35i15.17567>
- Liu Y, Li Q, Wang B, et al (2023) A survey of quantum-cognitively inspired sentiment analysis models. ACM Computing Surveys 56(1):1–37. <https://doi.org/10.1145/3604550>
- Lo KI, Sadrzadeh M, Mansfield S (2025) Quantum-like contextuality in large language models. Proceedings of the Royal Society A [arXiv:2412.16806](https://arxiv.org/abs/2412.16806)
- López-Incera A, Ried K, Müller T, et al (2020) Development of swarm behavior in artificial learning agents that adapt to different foraging environments. PLoS ONE 15(12):e0243628. <https://doi.org/10.1371/journal.pone.0243628>
- Lorenz R, Pearson A, Meichanetzidis K, et al (2023) QNLP in practice: Running compositional models of meaning on a quantum computer. Journal of Artificial Intelligence Research 76:1305–1342. <https://doi.org/10.1613/jair.1.14329>
- Luckcuck M, Farrell M, Dennis LA, et al (2019) Formal specification and verification of autonomous robotic systems: A survey. ACM Computing Surveys 52(5):100:1–100:41. <https://doi.org/10.1145/3342355>
- Lukac M, Perkowski M (2007) Quantum mechanical model of emotional robot behaviors. In: Proceedings of the 37th International Symposium on Multiple-Valued Logic (ISMVL 2007). IEEE, Oulu, Finland
- Maksimovic M, Maksymov IS (2025) Quantum-cognitive neural networks: Assessing confidence and uncertainty with human decision-making simulations. Big Data and Cognitive Computing 9(1):12. <https://doi.org/10.3390/bdce9010012>
- Maksymov IS (2025) Cognition in superposition: quantum models in AI, finance, defence, gaming and collective behaviour. [arXiv:2508.20098](https://arxiv.org/abs/2508.20098)
- Mannone M, Seidita V, Chella A (2022) Categories, quantum computing, and swarm robotics: A case study. Mathematics 10(3):372. <https://doi.org/10.3390/math10030372>
- Mannone M, Chella A, Pilato G, et al (2024) Modeling robotic thinking and creativity: A classic–quantum dialogue. Mathematics 12(5):642. <https://doi.org/10.3390/math12050642>
- Martínez-Martínez I, Sánchez-Burillo E (2016) Quantum stochastic walks on networks for decision-making. Scientific Reports 6:23812. <https://doi.org/10.1038/srep23812>

- van der Meer J, Hoyte P, Roeder L, et al (2025) Modeling the quantum-like dynamics of human reliability ratings in human–AI interactions by interaction dependent Hamiltonians. ArXiv preprint, [arXiv:2504.13918](https://arxiv.org/abs/2504.13918)
- Meghdadi A, Akbarzadeh-T. MR, Javidan K (2022) A quantum-like cognitive approach to modeling human biased selection behavior. *Scientific Reports* 12:22545. <https://doi.org/10.1038/s41598-022-13757-2>
- Melnikov AA, Makmal A, Briegel HJ (2018) Benchmarking projective simulation in navigation problems. *IEEE Access* 6:64639–64648. <https://doi.org/10.1109/ACCESS.2018.2876494>
- Meyer N, Ufrecht C, Periyasamy M, et al (2022) A survey on quantum reinforcement learning. [arXiv:2211.03464](https://arxiv.org/abs/2211.03464)
- Moreira C (2017) Quantum probabilistic graphical models for cognition and decision. PhD thesis, Instituto Superior Técnico, Universidade de Lisboa, Lisbon, Portugal
- Moreira C, Wichert A (2014) Interference effects in quantum belief networks. *Applied Soft Computing* 25:64–85. <https://doi.org/10.1016/j.asoc.2014.09.008>
- Moreira C, Wichert A (2016) Quantum-like Bayesian networks for modeling decision making. *Frontiers in Psychology* 7:11. <https://doi.org/10.3389/fpsyg.2016.00011>
- Moreira C, Wichert A (2017) Are quantum models for order effects quantum? *International Journal of Theoretical Physics* 56:4029–4046. <https://doi.org/10.1007/s10773-017-3424-5>
- Moreira C, Fell L, Dehdashti S, et al (2020a) Towards a quantum-like cognitive architecture for decision-making. *Behavioral and Brain Sciences* 43:e17. <https://doi.org/10.1017/S0140525X19001687>
- Moreira C, Tiwari P, Pandey HM, et al (2020b) Quantum-like influence diagrams for decision-making. *Neural Networks* 132:190–210. <https://doi.org/10.1016/j.neunet.2020.07.009>
- Nausheen F, Ahmed K, Khan MI, et al (2025) Quantum natural language processing: A comprehensive review of models, methods, and applications. [arXiv:2504.09909](https://arxiv.org/abs/2504.09909)
- Nebli A, Saadatdoorabi H, Yam K (2026) Deep sequence modeling with quantum dynamics: Language as a wave function. ArXiv preprint, [arXiv:2602.22255](https://arxiv.org/abs/2602.22255)
- von Neumann J (1932) *Mathematische Grundlagen der Quantenmechanik*. Springer, Berlin, english translation: *Mathematical Foundations of Quantum Mechanics*, trans. R. T. Beyer, Princeton University Press, Princeton, 1955
- Ozawa M, Khrennikov A (2020) Application of theory of quantum instruments to psychology: Combination of question order effect with response replicability effect.

- Entropy 22(1):37. <https://doi.org/10.3390/e22010037>
- Ozawa M, Khrennikov A (2021) Modeling combination of question order effect, response replicability effect, and QQ-equality with quantum instruments. *Journal of Mathematical Psychology* 100:102491. <https://doi.org/10.1016/j.jmp.2020.102491>
- Page MJ, McKenzie JE, Bossuyt PM, et al (2021) The PRISMA 2020 statement: an updated guideline for reporting systematic reviews. *BMJ* 372:n71. <https://doi.org/10.1136/bmj.n71>
- Paparo GD, Dunjko V, Makmal A, et al (2014) Quantum speedup for active learning agents. *Physical Review X* 4(3):031002. <https://doi.org/10.1103/PhysRevX.4.031002>
- Petschnigg C, Brandstötter M, Pichler H, et al (2019) Quantum computation in robotic science and applications. In: 2019 International Conference on Robotics and Automation (ICRA). IEEE, pp 803–810, <https://doi.org/10.1109/ICRA.2019.8793768>
- Pothos EM, Busemeyer JR (2009) A quantum probability explanation for violations of ‘rational’ decision theory. *Proceedings of the Royal Society B: Biological Sciences* 276(1665):2171–2178. <https://doi.org/10.1098/rspb.2009.0121>
- Pothos EM, Busemeyer JR (2013) Can quantum probability provide a new direction for cognitive modeling? *Behavioral and Brain Sciences* 36(3):255–274. <https://doi.org/10.1017/S0140525X12001525>
- Pothos EM, Busemeyer JR (2022) Quantum cognition. *Annual Review of Psychology* 73(1):749–778. <https://doi.org/10.1146/annurev-psych-033020-123501>
- Pothos EM, Busemeyer JR, Trueblood JS (2013) A quantum geometric model of similarity. *Psychological Review* 120(3):679–696. <https://doi.org/10.1037/a0033142>
- Pothos EM, Yearsley JM, Shiffrin RM, et al (2017) The rational status of quantum cognition. *Journal of Experimental Psychology: General* 146(7):968–987. <https://doi.org/10.1037/xge0000312>
- Preskill J (2018) Quantum computing in the NISQ era and beyond. *Quantum* 2:79. <https://doi.org/10.22331/q-2018-08-06-79>
- Pushparaj K, Ky G, Ayeni AJ, et al (2021) A quantum-inspired model for human-automation trust in air traffic controllers derived from functional magnetic resonance imaging and correlated with behavioural indicators. *Journal of Air Transport Management* 97:102143. <https://doi.org/10.1016/j.jairtraman.2021.102143>
- Ried K, Müller T, Briegel HJ (2019) Modelling collective motion based on the principle of agency: General framework and the case of marching locusts. *PLoS ONE*

- 14(2):e0212044. <https://doi.org/10.1371/journal.pone.0212044>
- Roeder L, Hoyte P, van der Meer J, et al (2023) A quantum model of trust calibration in human–AI interactions. *Entropy* 25(9):1362. <https://doi.org/10.3390/e25091362>
- Saggio V, Asenbeck BE, Hamann A, et al (2021) Experimental quantum speed-up in reinforcement learning agents. *Nature* 591(7849):229–233. <https://doi.org/10.1038/s41586-021-03242-7>
- Savage LJ (1954) *The Foundations of Statistics*. John Wiley & Sons, New York
- Schuld M, Killoran N (2019) Quantum machine learning in feature Hilbert spaces. *Physical Review Letters* 122(4):040504. <https://doi.org/10.1103/PhysRevLett.122.040504>
- Shafir E, Tversky A (1992) Thinking through uncertainty: Nonconsequential reasoning and choice. *Cognitive Psychology* 24(4):449–474. [https://doi.org/10.1016/0010-0285\(92\)90015-T](https://doi.org/10.1016/0010-0285(92)90015-T)
- She L, Han S, Liu X (2021) Application of quantum-like Bayesian network and belief entropy for interference effect in multi-attribute decision making problem. *Computers & Industrial Engineering* 157:107307. <https://doi.org/10.1016/j.cie.2021.107307>
- Sinha A, Macaluso A, Klusch M (2025) Nav-Q: Quantum deep reinforcement learning for collision-free navigation of self-driving cars. *Quantum Machine Intelligence* 7(1):19. <https://doi.org/10.1007/s42484-024-00226-4>
- Skolik A, Jerbi S, Dunjko V (2022) Quantum agents in the Gym: A variational quantum algorithm for deep Q-learning. *Quantum* 6:720. <https://doi.org/10.22331/q-2022-05-24-720>
- Song Q, Wang W, Fu W, et al (2022) Research on quantum cognition in autonomous driving. *Scientific Reports* 12:300. <https://doi.org/10.1038/s41598-021-04239-y>
- Sordoni A, Nie JY, Bengio Y (2013) Modeling term dependencies with quantum language models for IR. In: *Proceedings of the 36th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '13)*. ACM, Dublin, Ireland, pp 653–662, <https://doi.org/10.1145/2484028.2484098>
- Stoudenmire E, Schwab DJ (2016) Supervised learning with tensor networks. In: Lee DD, Sugiyama M, Luxburg UV, et al (eds) *Advances in Neural Information Processing Systems 29 (NIPS 2016)*. Curran Associates, Inc., pp 4799–4807, URL <https://proceedings.neurips.cc/paper/2016/hash/5314b9674c86e3f9d1ba25ef9bb32895-Abstract.html>, arXiv:1605.05775

- Sutton RS, Barto AG (2018) Reinforcement Learning: An Introduction, 2nd edn. MIT Press, Cambridge, MA
- Svozil K (2025) Quantum contextuality for contextual word embeddings. [arXiv:2504.13824](https://arxiv.org/abs/2504.13824)
- Tang E (2019) A quantum-inspired classical algorithm for recommendation systems. In: Proceedings of the 51st Annual ACM SIGACT Symposium on Theory of Computing (STOC 2019). ACM, New York, NY, USA, pp 217–228, <https://doi.org/10.1145/3313276.3316310>
- Tegmark M (2000) Importance of quantum decoherence in brain processes. *Physical Review E* 61(4):4194–4206. <https://doi.org/10.1103/PhysRevE.61.4194>
- Thrun S, Burgard W, Fox D (2005) Probabilistic Robotics. MIT Press, Cambridge, MA
- Tjuaatja L, Chen V, Wu T, et al (2024) Do LLMs exhibit human-like response biases? A case study in survey design. *Transactions of the Association for Computational Linguistics* 12:1011–1026. <https://doi.org/10.1162/tacl.a.00685>
- Trueblood JS, Busemeyer JR (2011) A quantum probability account of order effects in inference. *Cognitive Science* 35(8):1518–1552. <https://doi.org/10.1111/j.1551-6709.2011.01197.x>
- Trueblood JS, Hemmer P (2017) The generalized quantum episodic memory model. *Cognitive Science* 41(8):2089–2125. <https://doi.org/10.1111/cogs.12460>
- Trueblood JS, Yearsley JM, Pothos EM (2017) A quantum probability framework for human probabilistic inference. *Journal of Experimental Psychology: General* 146(9):1307–1341. <https://doi.org/10.1037/xge0000326>
- Tversky A, Kahneman D (1974) Judgment under uncertainty: Heuristics and biases. *Science* 185(4157):1124–1131. <https://doi.org/10.1126/science.185.4157.1124>
- Tversky A, Kahneman D (1983) Extensional versus intuitive reasoning: The conjunction fallacy in probability judgment. *Psychological Review* 90(4):293–315. <https://doi.org/10.1037/0033-295X.90.4.293>
- Upreti S, Gkoumas D, Song D (2020) A survey of quantum theory inspired approaches to information retrieval. *ACM Computing Surveys* 53(5):98:1–98:39. <https://doi.org/10.1145/3402179>
- Upreti S, Liu H, Jaiswal AK, et al (2024) Investigating context effects in similarity judgements in large language models. [arXiv:2408.10711](https://arxiv.org/abs/2408.10711)
- Veloz T, Sobetska O (2024) Analyzing the conjunction fallacy as a fact. [arXiv:2402.13615](https://arxiv.org/abs/2402.13615)

- Waddup O, Yearsley JM, Blasiak P, et al (2023) Temporal Bell inequalities in cognition. *Psychonomic Bulletin & Review* 30(5):1946–1953. <https://doi.org/10.3758/s13423-023-02275-5>
- Waddup OJ, Blasiak P, Yearsley JM, et al (2021) Sensitivity to context in human interactions. *Mathematics* 9(21):2784. <https://doi.org/10.3390/math9212784>
- Wang Z, Busemeyer JR (2013) A quantum question order model supported by empirical tests of an a priori and precise prediction. *Topics in Cognitive Science* 5(4):689–710. <https://doi.org/10.1111/tops.12040>
- Wang Z, Busemeyer JR (2016) Interference effects of categorization on decision making. *Cognition* 150:133–149. <https://doi.org/10.1016/j.cognition.2016.01.019>
- Wang Z, Busemeyer JR, Atmanspacher H, et al (2013) The potential of using quantum theory to build models of cognition. *Topics in Cognitive Science* 5(4):672–688. <https://doi.org/10.1111/tops.12043>
- Wang Z, Solloway T, Shiffrin RM, et al (2014) Context effects produced by question orders reveal quantum nature of human judgments. *Proceedings of the National Academy of Sciences* 111(26):9431–9436. <https://doi.org/10.1073/pnas.1407756111>
- Wei Q, Ma H, Chen C, et al (2022) Deep reinforcement learning with quantum-inspired experience replay. *IEEE Transactions on Cybernetics* 52(9):9326–9338. <https://doi.org/10.1109/TCYB.2021.3053414>
- Wichert A, Moreira C, Bruza P (2020) Balanced quantum-like Bayesian networks. *Entropy* 22(2):170. <https://doi.org/10.3390/e22020170>
- Widdows D, Peters S (2003) Word vectors and quantum logic: Experiments with negation and disjunction. In: *Proceedings of the 8th Mathematics of Language Conference (MOL 8)*, Bloomington, IN, pp 141–154
- Widdows D, Kitto K, Cohen T (2021) Quantum mathematics in artificial intelligence. *Journal of Artificial Intelligence Research* 72:1307–1341. <https://doi.org/10.1613/jair.1.12702>
- Widdows D, Rani J, Pothos EM (2023) Quantum circuit components for cognitive decision-making. *Entropy* 25(4):548. <https://doi.org/10.3390/e25040548>
- Widdows D, Aboumradi W, Kim D, et al (2024) Quantum natural language processing. *KI – Künstliche Intelligenz* 38(4):293–310. <https://doi.org/10.1007/s13218-024-00861-w>
- Yan F, Iliyasu AM, Jiao S, et al (2019) Quantum structure for modelling emotion space of robots. *Applied Sciences* 9(16):3351. <https://doi.org/10.3390/app9163351>

- Yan F, Iliyasa AM, Hirota K (2021a) Conceptual framework for quantum affective computing and its use in fusion of multi-robot emotions. *Electronics* 10(2):100. <https://doi.org/10.3390/electronics10020100>
- Yan F, Iliyasa AM, Hirota K (2021b) Emotion space modelling for social robots. *Engineering Applications of Artificial Intelligence* 100:104178. <https://doi.org/10.1016/j.engappai.2021.104178>
- Yan F, Yang X, Li N, et al (2021c) Emotion generation and transition of companion robots based on Plutchik's model and quantum circuit schemes. *Security and Communication Networks* 2021:6802521. <https://doi.org/10.1155/2021/6802521>
- Yan F, Iliyasa AM, Li N, et al (2024) Quantum robotics: a review of emerging trends. *Quantum Machine Intelligence* 6(2):86. <https://doi.org/10.1007/s42484-024-00225-5>
- Yearsley JM (2017) Advanced tools and concepts for quantum cognition: A tutorial. *Journal of Mathematical Psychology* 78:24–39. <https://doi.org/10.1016/j.jmp.2016.07.005>
- Yearsley JM, Busemeyer JR (2016) Quantum cognition and decision theories: A tutorial. *Journal of Mathematical Psychology* 74:99–116. <https://doi.org/10.1016/j.jmp.2015.11.005>
- Yearsley JM, Pothos EM (2016) Zeno's paradox in decision-making. *Proceedings of the Royal Society B: Biological Sciences* 283(1828):20160291. <https://doi.org/10.1098/rspb.2016.0291>
- Yukalov VI (2020) Evolutionary processes in quantum decision theory. *Entropy* 22(6):681. <https://doi.org/10.3390/e22060681>
- Yukalov VI (2023) Quantum operation of affective artificial intelligence. *Laser Physics* 33(6):065204. <https://doi.org/10.1088/1555-6611/accf7a>
- Yukalov VI, Sornette D (2011) Decision theory with prospect interference and entanglement. *Theory and Decision* 70(3):283–328. <https://doi.org/10.1007/s11238-010-9202-y>
- Yukalov VI, Sornette D (2016) Quantum probability and quantum decision-making. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences* 374(2058):20150100. <https://doi.org/10.1098/rsta.2015.0100>, [arXiv:1601.02480](https://arxiv.org/abs/1601.02480)
- Yukalov VI, Yukalova EP, Sornette D (2018) Information processing by networks of quantum decision makers. *Physica A: Statistical Mechanics and its Applications* 492:747–766. <https://doi.org/10.1016/j.physa.2017.11.004>

- Yukalov VI, Yukalova EP, Sornette D (2022) Role of collective information in networks of quantum operating agents. *Physica A: Statistical Mechanics and its Applications* 598:127365. <https://doi.org/10.1016/j.physa.2022.127365>
- Yun WJ, Kwak Y, Kim JP, et al (2022) Quantum multi-agent reinforcement learning via variational quantum circuit design. In: *Proceedings of the 2022 IEEE 42nd International Conference on Distributed Computing Systems (ICDCS)*. IEEE, pp 1332–1335, <https://doi.org/10.1109/ICDCS54860.2022.00151>
- Zhang C, Li Q, Song D, et al (2025) Quantum-inspired semantic matching based on neural networks with the duality of density matrices. *Engineering Applications of Artificial Intelligence* 140:109667. <https://doi.org/10.1016/j.engappai.2024.109667>
- Zhang P, Su Z, Zhang L, et al (2018) A quantum many-body wave function inspired language modeling approach. In: *Proceedings of the 27th ACM International Conference on Information and Knowledge Management (CIKM 2018)*. ACM, pp 1303–1312, <https://doi.org/10.1145/3269206.3271723>